

IF2230 Threads



Threads

- ▶ Overview
- ▶ Multicore Programming
- ▶ Multithreading Models
- ▶ Thread Libraries
- ▶ Threading Issues
- ▶ Examples



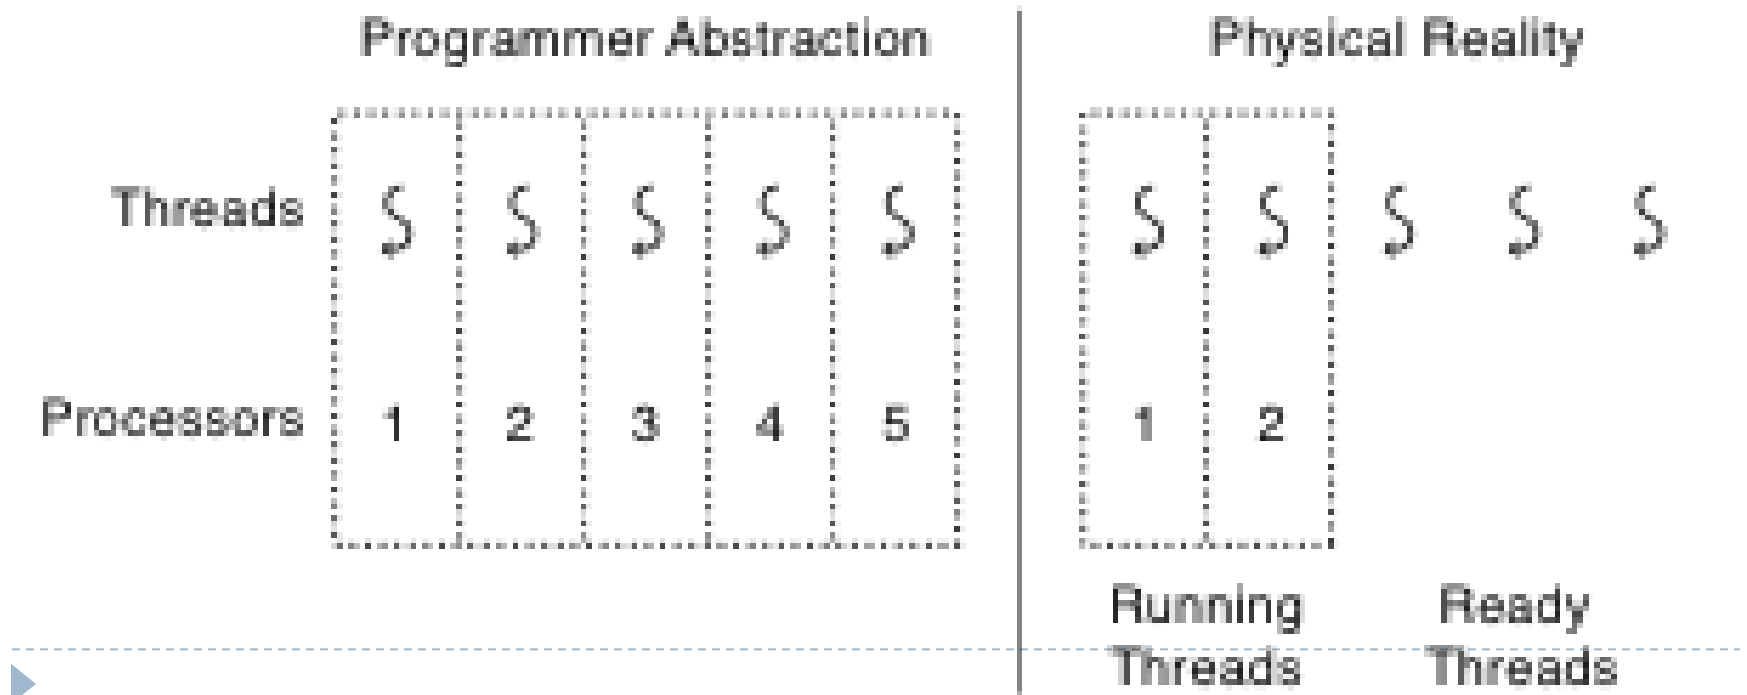
Definisi

- ▶ **Thread: sekuens eksekusi tunggal yang merepresentasikan task yang dapat dijadwalkan tersendiri**
 - ▶ sekuens eksekusi tunggal: model pemrograman yang sederhana
 - ▶ dapat dijadwalkan tersendiri: OS dapat menjalankan atau men-suspend sebuah thread kapan saja
- ▶ **Proteksi (akses memori, resources)**
 - ▶ 1 atau beberapa thread dapat menshare domain proteksi yang sama



Abstraksi Thread

- ▶ Jumlah prosesor yang tak terbatas
- ▶ Threads dapat berjalan dengan kecepatan yang bervariasi
 - ▶ Program harus dirancang agar dapat bekerja dengan berbagai kemungkinan penjadwalan




Programmer vs. Processor View

Programmer's View	Possible Execution #1	Possible Execution #2	Possible Execution #3
.	.	.	.
.	.	.	.
.	.	.	.
$x = x + 1;$	$x = x + 1;$	$x = x + 1;$	$x = x + 1;$
$y = y * x;$	$y = y * x;$	$y = y * x;$
$z = x * 5y;$	$z = x * 5y;$	Thread is suspended. Other thread(s) run. Thread is resumed. Thread is suspended. Other thread(s) run. Thread is resumed.
.
.	.	$y = y * x;$	$z = x * 5y;$
.	.	$z = x * 5y;$	




Possible Executions


One Execution



Another Execution



Another Execution



Thread Operations

- ▶ `thread_create(thread, func, args)`
 - ▶ Create a new thread to run `func(args)`
- ▶ `thread_yield()`
 - ▶ Relinquish processor voluntarily
- ▶ `thread_join(thread)`
 - ▶ In parent, wait for forked thread to exit, then return
- ▶ `thread_exit`
 - ▶ Quit thread and clean up, wake up joiner if any



Example: threadHello

```
#define NTHREADS 10
thread_t threads[NTHREADS];
main() {
    for (i = 0; i < NTHREADS; i++) thread_create(&threads[i], &go, i);
    for (i = 0; i < NTHREADS; i++) {
        exitValue = thread_join(threads[i]);
        printf("Thread %d returned with %ld\n", i, exitValue);
    }
    printf("Main thread done.\n");
}

void go (int n) {
    printf("Hello from thread %d\n", n);
    thread_exit(100 + n);
    // REACHED?
}
```

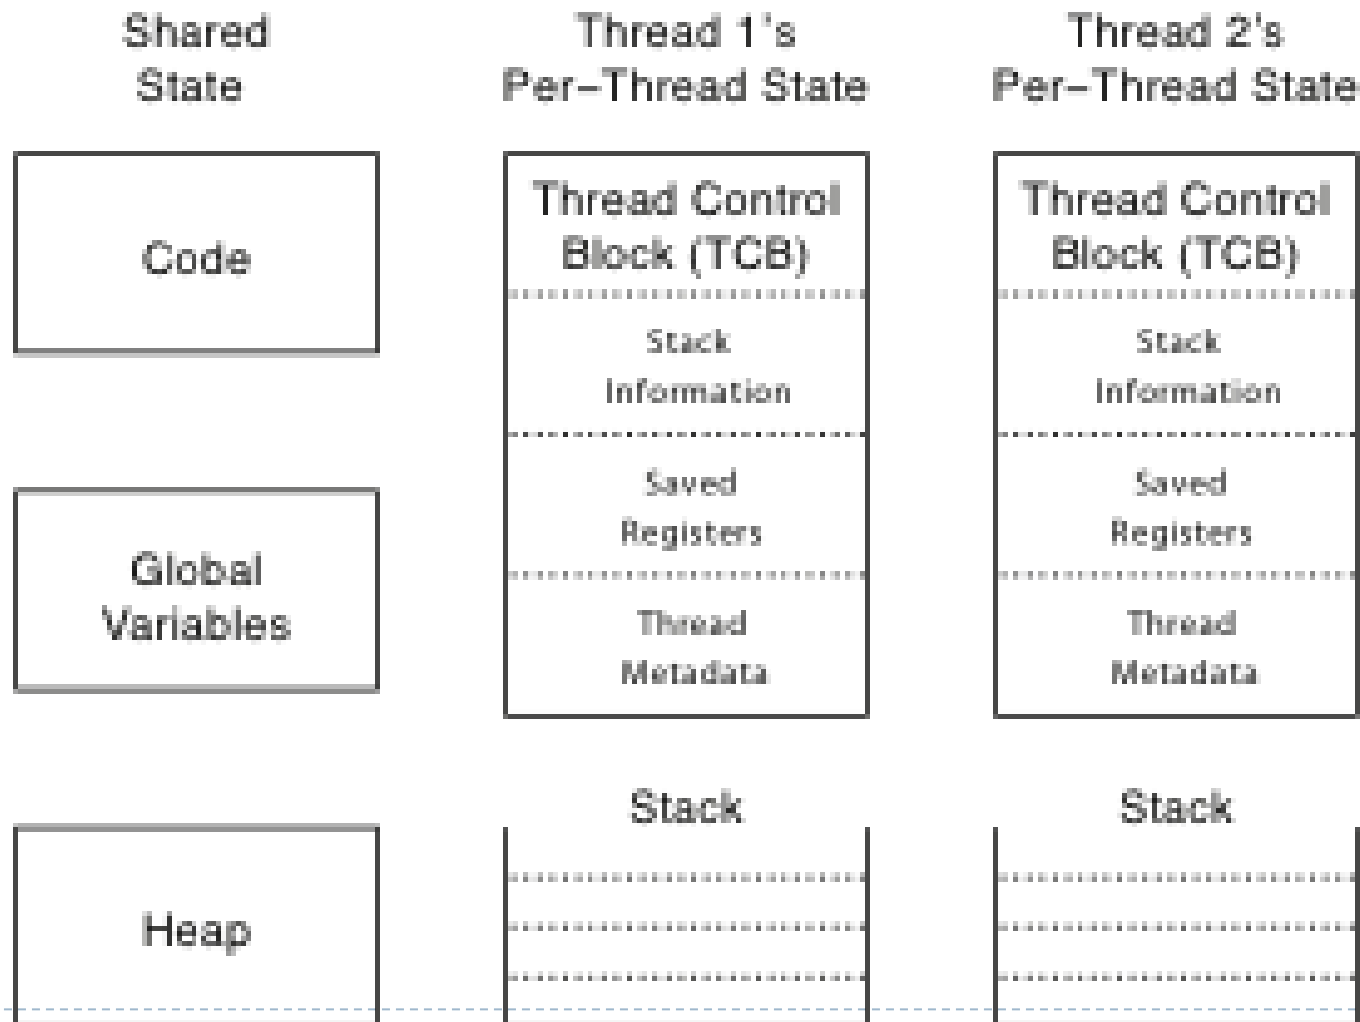
threadHello: Example Output

- ▶ Why must “thread returned” print in order?
- ▶ What is maximum # of threads running when thread 5 prints hello?
- ▶ Minimum?

```
bash-3.2$ ./threadHello
Hello from thread 0
Hello from thread 1
Thread 0 returned 100
Hello from thread 3
Hello from thread 4
Thread 1 returned 101
Hello from thread 5
Hello from thread 2
Hello from thread 6
Hello from thread 8
Hello from thread 7
Hello from thread 9
Thread 2 returned 102
Thread 3 returned 103
Thread 4 returned 104
Thread 5 returned 105
Thread 6 returned 106
Thread 7 returned 107
Thread 8 returned 108
Thread 9 returned 109
Main thread done.
```



Thread Data Structures

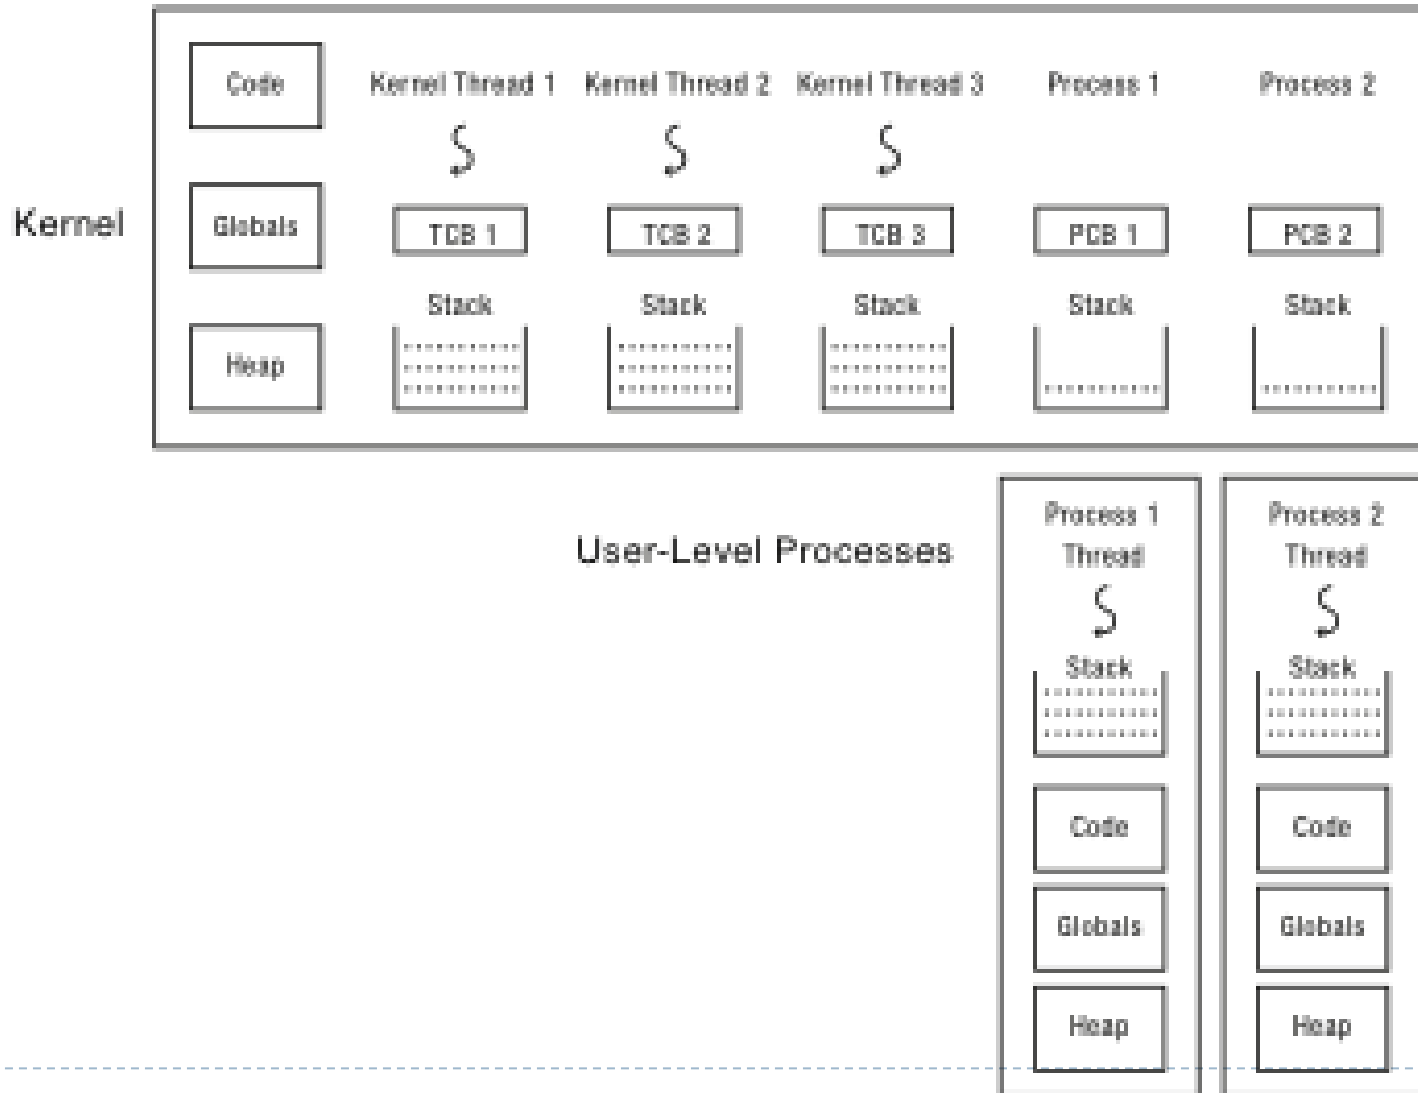


Implementing Threads: Roadmap

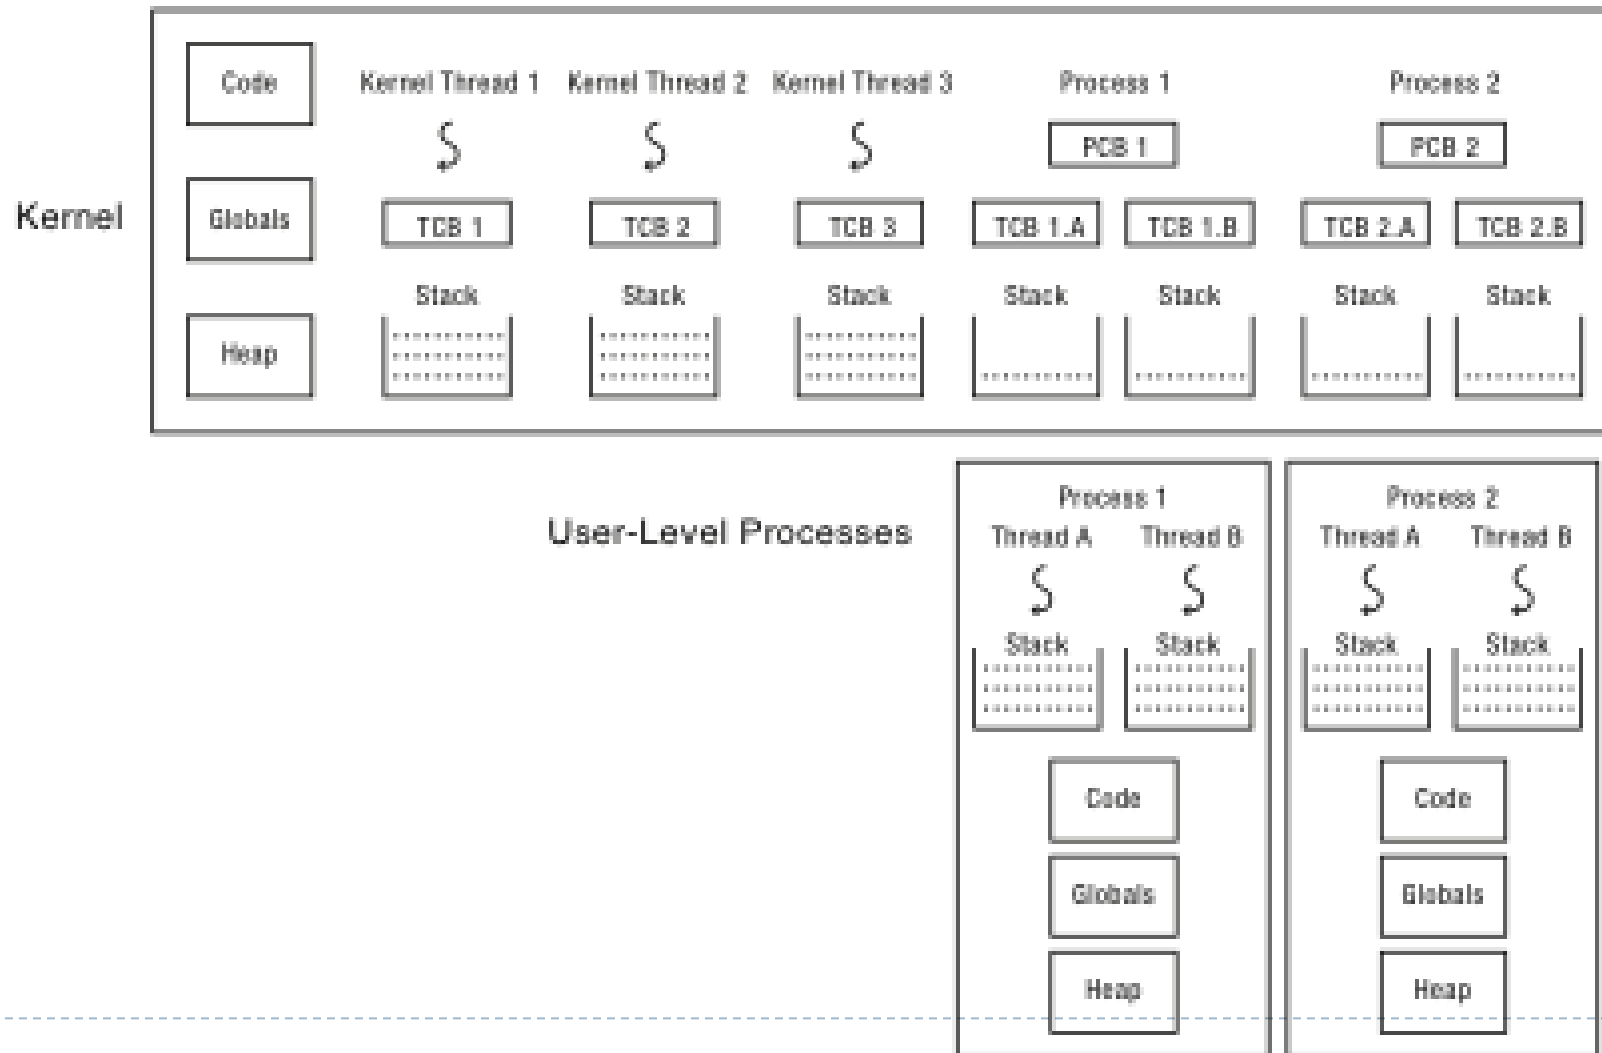
- ▶ **Kernel threads**
 - ▶ Thread abstraction only available to kernel
 - ▶ To the kernel, a kernel thread and a single threaded user process look quite similar
- ▶ **Multithreaded processes using kernel threads (Linux, MacOS)**
 - ▶ Kernel thread operations available via syscall
- ▶ **User-level threads**
 - ▶ Thread operations without system calls



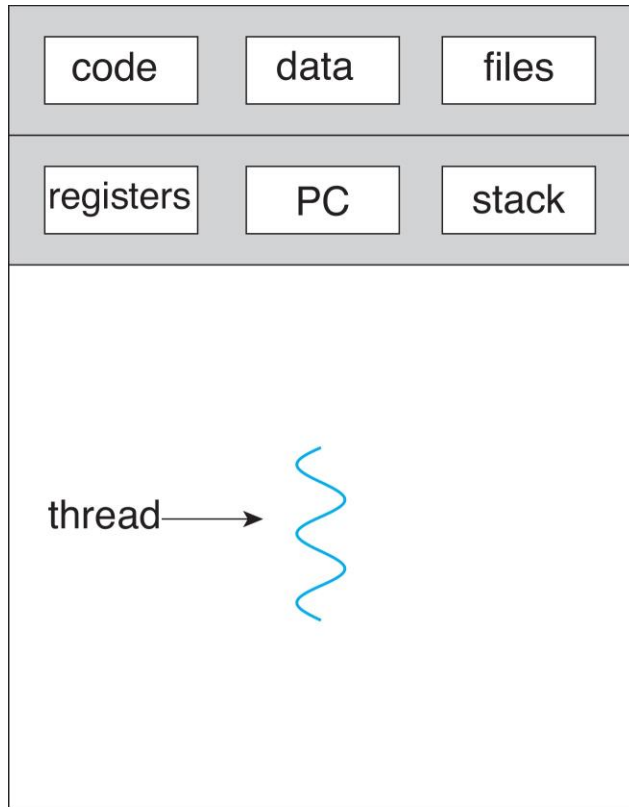
Multithreaded OS Kernel




Multithreaded User Processes (Take 1)



Single and Multithreaded Processes

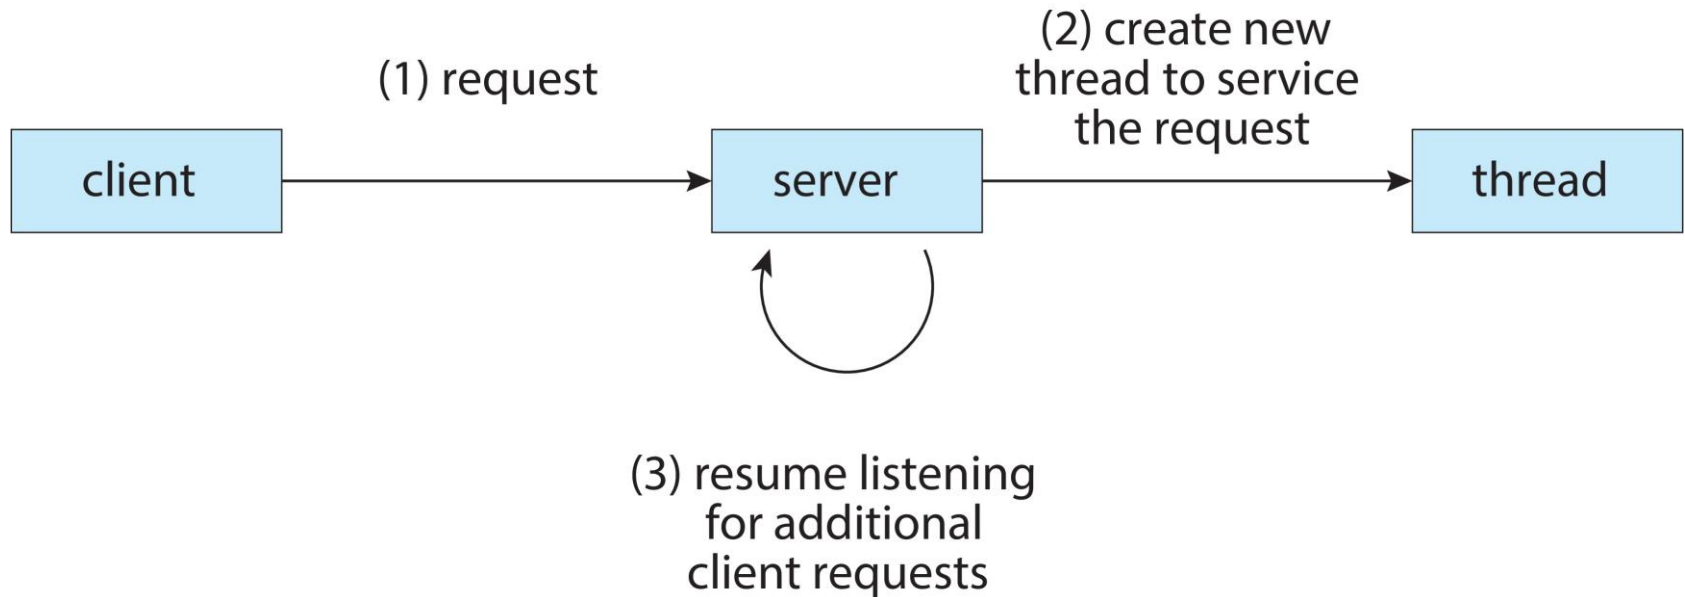


single-threaded process



multithreaded process

Multithreaded Server Architecture



Benefits

- ▶ **Servers**
 - ▶ Multiple connections handled simultaneously
- ▶ **Parallel programs**
 - ▶ To achieve better performance
- ▶ **Programs with user interfaces**
 - ▶ To achieve user responsiveness while doing computation
- ▶ **Network and disk bound programs**
 - ▶ To hide network/disk latency




Multicore Programming

- ▶ **Multicore** or **multiprocessor** systems puts pressure on programmers, challenges include:
 - ▶ **Dividing activities**
 - ▶ **Balance**
 - ▶ **Data splitting**
 - ▶ **Data dependency**
 - ▶ **Testing and debugging**
- ▶ **Parallelism** implies a system can perform more than one task simultaneously
- ▶ **Concurrency** supports more than one task making progress
 - ▶ Single processor / core, scheduler providing concurrency




Concurrency vs. Parallelism

- **Concurrent execution on single-core system:**



- **Parallelism on a multi-core system:**



User Threads

- ▶ Thread management done by user-level threads library
- ▶ Three primary thread libraries:
 - ▶ POSIX Pthreads
 - ▶ Win32 threads
 - ▶ Java threads



Kernel Threads


- ▶ Supported by the Kernel

- ▶ Examples

- ▶ Windows XP/2000
- ▶ Solaris
- ▶ Linux
- ▶ Tru64 UNIX
- ▶ Mac OS X



User and Kernel Threads




Multithreading Models

- ▶ Many-to-One
- ▶ One-to-One
- ▶ Many-to-Many




Many-to-One

- ▶ Many user-level threads mapped to single kernel thread
- ▶ One thread blocking causes all to block
- ▶ Multiple thread may not run in parallel on multicore system
- ▶ Examples:
 - ▶ Solaris Green Threads
 - ▶ GNU Portable Threads




One-to-One

- ▶ Each user-level thread maps to kernel thread
- ▶ Creating a user-level thread creates a kernel thread
- ▶ More concurrency than many-to-one
- ▶ Examples
 - ▶ Windows NT/XP/2000
 - ▶ Linux
 - ▶ Solaris 9 and later




Many-to-Many Model

- ▶ Allows many user level threads to be mapped to many kernel threads
- ▶ Allows the operating system to create a sufficient number of kernel threads
- ▶ Solaris prior to version 9
- ▶ Windows NT/2000 with the *ThreadFiber* package



Two-level Model

- ▶ Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- ▶ Examples
 - ▶ IRIX
 - ▶ HP-UX
 - ▶ Tru64 UNIX
 - ▶ Solaris 8 and earlier



Thread Libraries

- ▶ **Thread library** provides programmer with API for creating and managing threads
- ▶ Two primary ways of implementing
 - ▶ Library entirely in user space
 - ▶ Kernel-level library supported by the OS



Pthreads

- ▶ May be provided either as user-level or kernel-level
- ▶ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ▶ ***Specification***, not ***implementation***
- ▶ API specifies behavior of the thread library, implementation is up to development of the library
- ▶ Common in UNIX operating systems (Linux & Mac OS X)



Pthreads Example

```
-----  
#include <pthread.h>  
#include <stdio.h>  
  
#include <stdlib.h>  
  
int sum; /* this data is shared by the thread(s) */  
void *runner(void *param); /* threads call this function */  
  
int main(int argc, char *argv[])  
{  
    pthread_t tid; /* the thread identifier */  
    pthread_attr_t attr; /* set of thread attributes */  
  
    /* set the default attributes of the thread */  
    pthread_attr_init(&attr);  
    /* create the thread */  
    pthread_create(&tid, &attr, runner, argv[1]);  
    /* wait for the thread to exit */  
    pthread_join(tid, NULL);  
  
    printf("sum = %d\n", sum);  
}
```



Pthreads Example (Cont.)

```
/* The thread will execute in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```



Pthreads Code for Joining 10 Threads

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```



Windows Multithreaded C Program

```
#include <windows.h>
#include <stdio.h>
DWORD Sum; /* data is shared by the thread(s) */

/* The thread will execute in this function */
DWORD WINAPI Summation(LPVOID Param)
{
    DWORD Upper = *(DWORD*)Param;
    for (DWORD i = 1; i <= Upper; i++)
        Sum += i;
    return 0;
}
```



Windows Multithreaded C Program (Cont.)

```
-----
int main(int argc, char *argv[])
{
    DWORD ThreadId;
    HANDLE ThreadHandle;
    int Param;

    Param = atoi(argv[1]);
    /* create the thread */
    ThreadHandle = CreateThread(
        NULL, /* default security attributes */
        0, /* default stack size */
        Summation, /* thread function */
        &Param, /* parameter to thread function */
        0, /* default creation flags */
        &ThreadId); /* returns the thread identifier */

    /* now wait for the thread to finish */
    WaitForSingleObject(ThreadHandle,INFINITE);

    /* close the thread handle */
    CloseHandle(ThreadHandle);

    printf("sum = %d\n",Sum);
}
-----
```



Java Threads

- ▶ Java threads are managed by the JVM
- ▶ Typically implemented using the threads model provided by underlying OS
- ▶ Java threads may be created by:
 - ▶ Extending Thread class
 - ▶ Implementing the Runnable interface

```
public interface Runnable
{
    public abstract void run();
}
```

- ▶ Standard practice is to implement Runnable interface



Java Threads

Implementing Runnable interface:

```
class Task implements Runnable
{
    public void run() {
        System.out.println("I am a thread.");
    }
}
```

Creating a thread:

```
Thread worker = new Thread(new Task());
worker.start();
```

Waiting on a thread:

```
try {
    worker.join();
}
catch (InterruptedException ie) { }
```



Java Executor Framework

- ▶ Rather than explicitly creating threads, Java also allows thread creation around the Executor interface:

```
public interface Executor
{
    void execute(Runnable command);
}
```

- ▶ The Executor is used as follows:

```
Executor service = new Executor();
service.execute(new Task());
```



Java Executor Framework

```
import java.util.concurrent.*;

class Summation implements Callable<Integer>
{
    private int upper;
    public Summation(int upper) {
        this.upper = upper;
    }

    /* The thread will execute in this method */
    public Integer call() {
        int sum = 0;
        for (int i = 1; i <= upper; i++)
            sum += i;

        return new Integer(sum);
    }
}
```



Java Executor Framework (Cont.)

```
public class Driver
{
    public static void main(String[] args) {
        int upper = Integer.parseInt(args[0]);

        ExecutorService pool = Executors.newSingleThreadExecutor();
        Future<Integer> result = pool.submit(new Summation(upper));

        try {
            System.out.println("sum = " + result.get());
        } catch (InterruptedException | ExecutionException ie) { }
    }
}
```



Threading Issues

- ▶ Semantics of **fork()** and **exec()** system calls
- ▶ Thread cancellation
- ▶ Signal handling
- ▶ Thread pools
- ▶ Thread specific data
- ▶ Scheduler activations



Semantics of `fork()` and `exec()`

- ▶ Does **`fork()`** duplicate only the calling thread or all threads?
 - ▶ Some UNIXes have two versions of `fork`
- ▶ **`exec()`** usually works as normal – replace the running process including all threads



Thread Cancellation

- ▶ Terminating a thread before it has finished
- ▶ Thread to be canceled is **target thread**
- ▶ Three general approaches:
 - ▶ **Asynchronous cancellation** terminates the target thread immediately
 - ▶ **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
 - ▶ Not cancellable



Thread Cancellation (Cont.)

- ▶ Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state

Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- ▶ If thread has cancellation disabled, cancellation remains pending until thread enables it
- ▶ Default type is deferred
 - ▶ Cancellation only occurs when thread reaches **cancellation point**
 - ▶ i.e., `pthread_testcancel()`
 - ▶ Then **cleanup handler** is invoked
- ▶ On Linux systems, thread cancellation is handled through signals

Signal Handling

- ▶ Signals are used in UNIX systems to notify a process that a particular event has occurred
- ▶ A **signal handler** is used to process signals
 1. Signal is generated by particular event
 2. Signal is delivered to a process
 3. Signal is handled
- ▶ Options:
 - ▶ Deliver the signal to the thread to which the signal applies
 - ▶ Deliver the signal to every thread in the process
 - ▶ Deliver the signal to certain threads in the process
 - ▶ Assign a specific thread to receive all signals for the process



Thread Pools

- ▶ Create a number of threads in a pool where they await work
- ▶ Advantages:
 - ▶ Usually slightly faster to service a request with an existing thread than create a new thread
 - ▶ Allows the number of threads in the application(s) to be bound to the size of the pool




Thread Local Storage

- ▶ Thread Local Storage (TLS) allows each thread to have its own copy of data
- ▶ Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- ▶ Different from local variables
 - ▶ Local variables visible only during single function invocation
 - ▶ TLS visible across function invocations
- ▶ Similar to **static** data
 - ▶ TLS is unique to each thread



Scheduler Activations

- ▶ Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ▶ Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
 - ▶ Appears to be a virtual processor on which process can schedule user thread to run
 - ▶ Each LWP attached to kernel thread
 - ▶ How many LWPs to create?
- ▶ Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library
- ▶ This communication allows an application to maintain the correct number kernel threads



Windows Threads

- ▶ Windows API – primary API for Windows applications
- ▶ Implements the one-to-one mapping, kernel-level
- ▶ Each thread contains
 - ▶ A thread id
 - ▶ Register set representing state of processor
 - ▶ Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - ▶ Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- ▶ The register set, stacks, and private storage area are known as the **context** of the thread

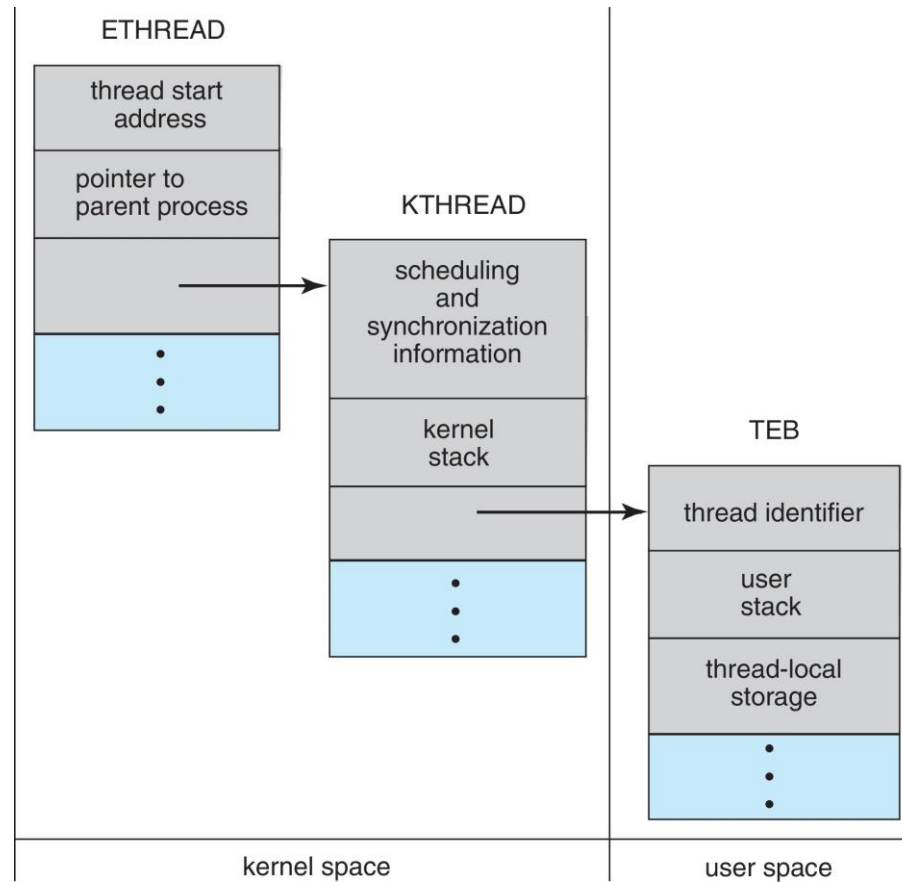


Windows Threads (Cont.)

- ▶ The primary data structures of a thread include:
 - ▶ ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space
 - ▶ KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
 - ▶ TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space



Windows Threads Data Structures



Linux Threads

- ▶ Linux refers to them as **tasks** rather than **threads**
- ▶ Thread creation is done through **clone()** system call
- ▶ **clone()** allows a child task to share the address space of the parent task (process)
 - ▶ Flags control behavior

flag	meaning
CLONE_FS	File-system information is shared.
CLONE_VM	The same memory space is shared.
CLONE_SIGHAND	Signal handlers are shared.
CLONE_FILES	The set of open files is shared.

- ▶ **struct task_struct** points to process data structures (shared or unique)