

IF2230

Introduction to Operating Systems

Intro OS Sem 2 – 2022-2023


Achmad Imam Kistijantoro (imam@staff.stei.itb.ac.id)

Judhi Santoso (judhi@staff.stei.itb.ac.id)

Robithoh Annur (robithoh@staff.stei.itb.ac.id)

What is an operating system?

- ▶ What is an *Operating System*?
 - ▶ The software interface between hardware and its users
- ▶ Operating systems:
 - ▶ Execute user and system programs
 - ▶ Manage and coordinate computer hardware
 - ▶ Serve as resource allocators
 - ▶ Are typically interrupt-driven

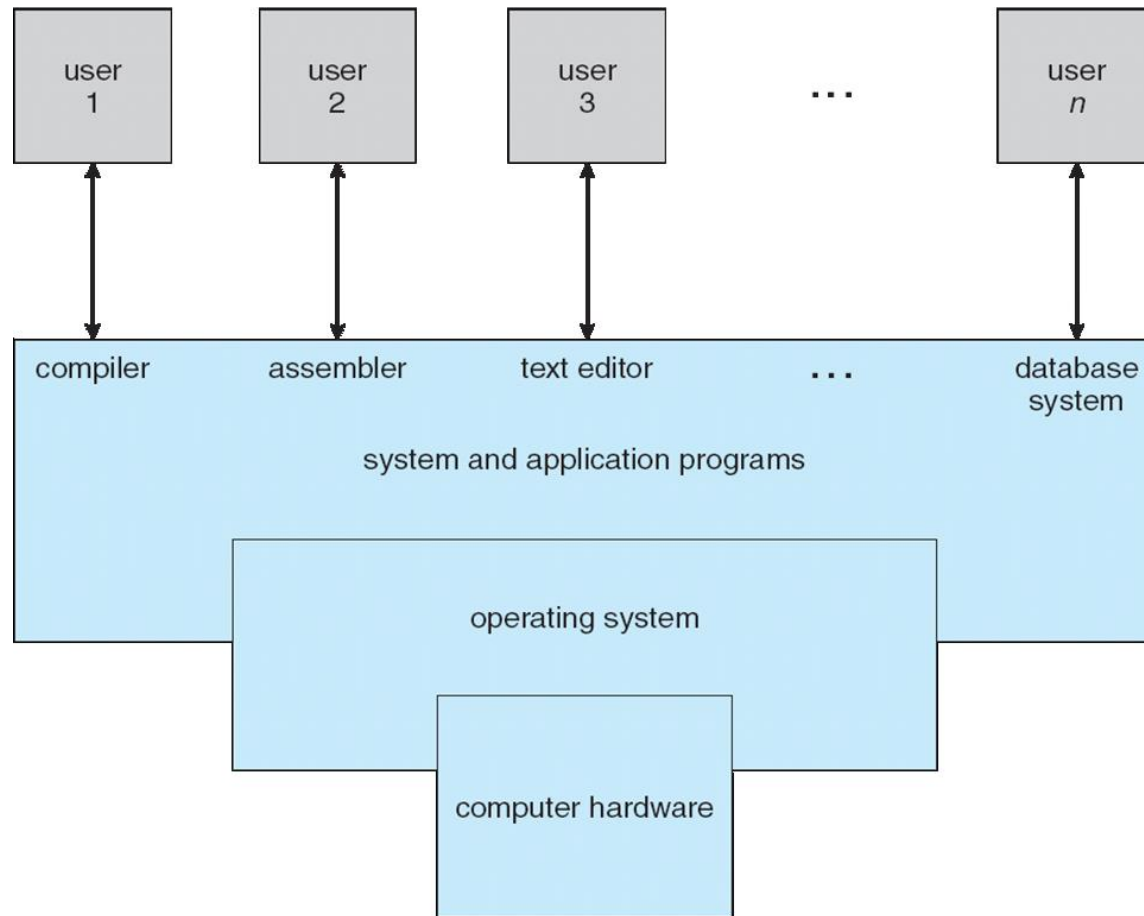


Computer System Structure

- ▶ Computer system can be divided into four components:
 - ▶ Hardware – provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - ▶ Operating system
 - ▶ Controls and coordinates use of hardware among various applications and users
 - ▶ Application programs – define the ways in which the system resources are used to solve the computing problems of the users
 - ▶ Word processors, compilers, web browsers, database systems, video games
 - ▶ Users
 - ▶ People, machines, other computers



Four Components of a Computer System



What Operating Systems Do

- ▶ Depends on the point of view
- ▶ Users want convenience, **ease of use**
 - ▶ Don't care about **resource utilization**
- ▶ But shared computer such as **mainframe** or **minicomputer** must keep all users happy
- ▶ Users of dedicate systems such as **workstations** have dedicated resources but frequently use shared resources from **servers**
- ▶ Handheld computers are resource poor, optimized for usability and battery life
- ▶ Some computers have little or no user interface, such as embedded computers in devices and automobiles



Operating System Definition

- ▶ OS is a **resource allocator**
 - ▶ Manages all resources
 - ▶ Decides between conflicting requests for efficient and fair resource use
- ▶ OS is a **control program**
 - ▶ Controls execution of programs to prevent errors and improper use of the computer



Operating System Definition (Cont.)

- ▶ No universally accepted definition
- ▶ “Everything a vendor ships when you order an operating system” is good approximation
 - ▶ But varies wildly
- ▶ “The one program running at all times on the computer” is the **kernel**. Everything else is either a system program (ships with the operating system) or an application program.



System software

- ▶ What is system software?
 - ▶ Computer programs that directly *control* the operations of the computer and its devices
- ▶ Operating systems:
 - ▶ Coordinate and orchestrate all activities of the hardware devices in a computer
 - ▶ Provide both a *Graphical User Interface (GUI)* and a *Command-Line Interface (CLI)* for its users



Operating system design goals

▶ From a user's perspective:

- ▶ easy to use
- ▶ easy to learn
- ▶ reliable
- ▶ safe
- ▶ fast
- ▶ etc.



■ System goals:

- reliability
- flexibility
- extensibility
- speed(y)
- efficiency
- maintainability
- etc.



Operating system services (i)

- ▶ An operating system provides services:
 - ▶ Program execution
 - ▶ Load programs into memory, run/suspend/halt programs, handle/display errors
 - ▶ I/O operations
 - ▶ Seamlessly interact with I/O devices, including disks, networks connection, etc.
 - ▶ Filesystem manipulation
 - ▶ Read/write/traverse filesystem directories, read/write files, enforce permissions, search for files




Operating system services (ii)

- ▶ Other operating system services:
 - ▶ Inter-Process Communications (IPC)
 - ▶ Processes exchange information via shared memory, message passing, sockets, pipes, files, etc.
 - ▶ Often spans multiple computers and networks
 - ▶ Error detection and recovery
 - ▶ Detect errors in CPU, memory, I/O devices, processes, network connections, etc.
 - ▶ Recover from errors gracefully, ensuring correct and consistent operations




History of operating systems (i)

- In the beginning...
...the 1940s




History of operating systems (ii)

► Automation in the 1950s with punch cards



Batch jobs


- ▶ A *job* is a unit of work submitted by a user to the operating system
- ▶ Jobs typically consist of:
 - ▶ a program either in a source language or in “executable” binary form
 - ▶ input data used by the program when it executes



History of operating systems (iii)

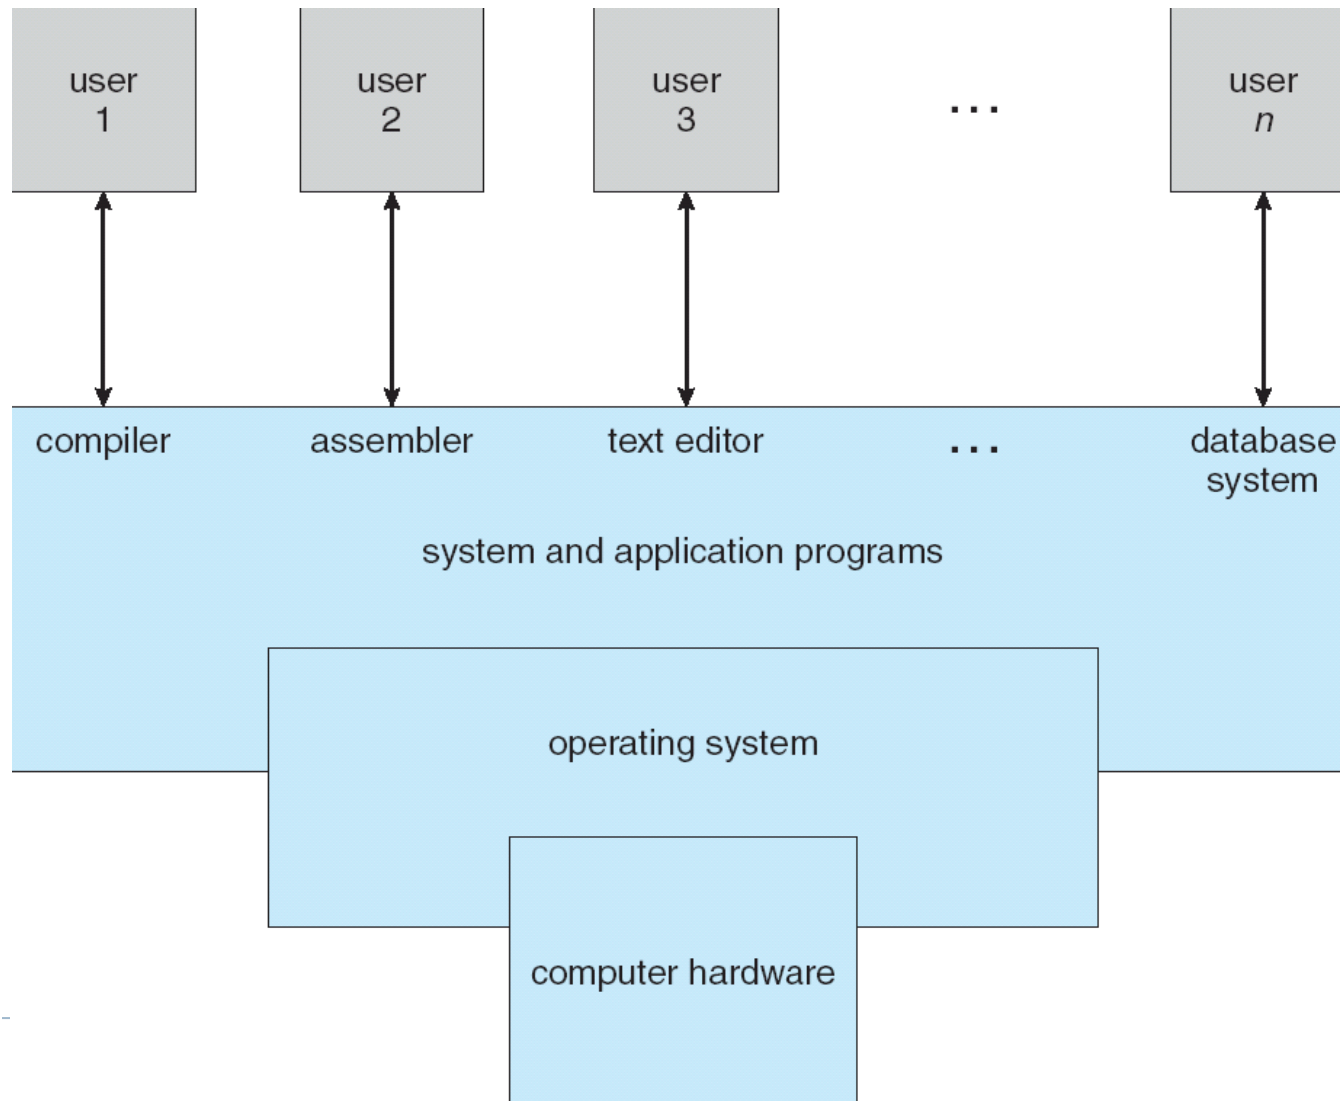
- ▶ IBM 360 introduced
(in 1964)

a computing
revolution!



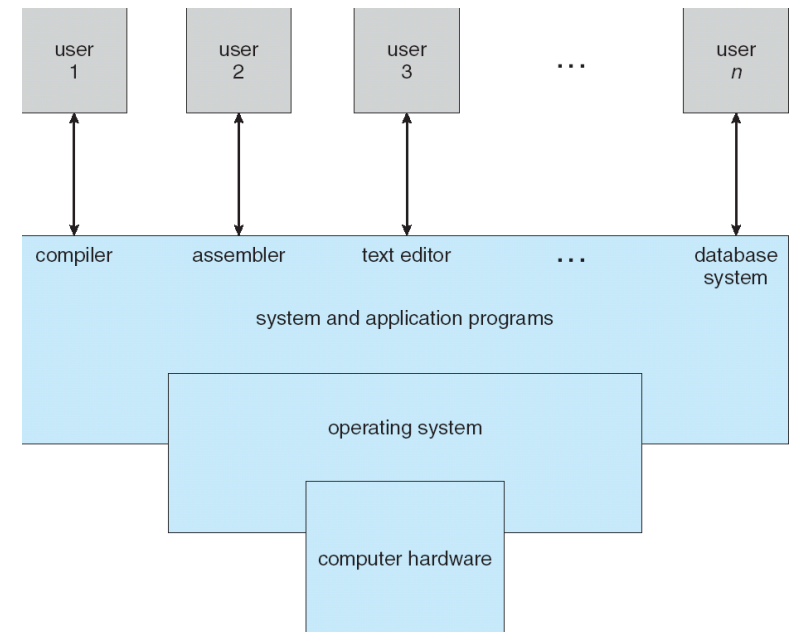
what's so
revolutionary?

Multiprogramming (i)



Multiprogramming (ii)

- ▶ In multiprogramming, several jobs reside in memory simultaneously
 - ▶ CPU use is shared and managed by the operating system




Multiprogramming (iii)

- ▶ Multiprogramming provides efficient use of the computer (CPU) and its resources (I/O)
 - ▶ One user cannot keep the CPU and I/O devices busy at all times
 - ▶ Multiprogramming organizes jobs such that the CPU always has exactly one to execute



Multiprogramming (iv)

- ▶ Computer is often idle – why?
 - ▶ CPU and hardware *significantly* faster than I/O
 - ▶ When a user or process is blocked waiting for I/O, the operating system switches to another job
 - ▶ A subset of jobs is stored in memory, awaiting CPU or I/O



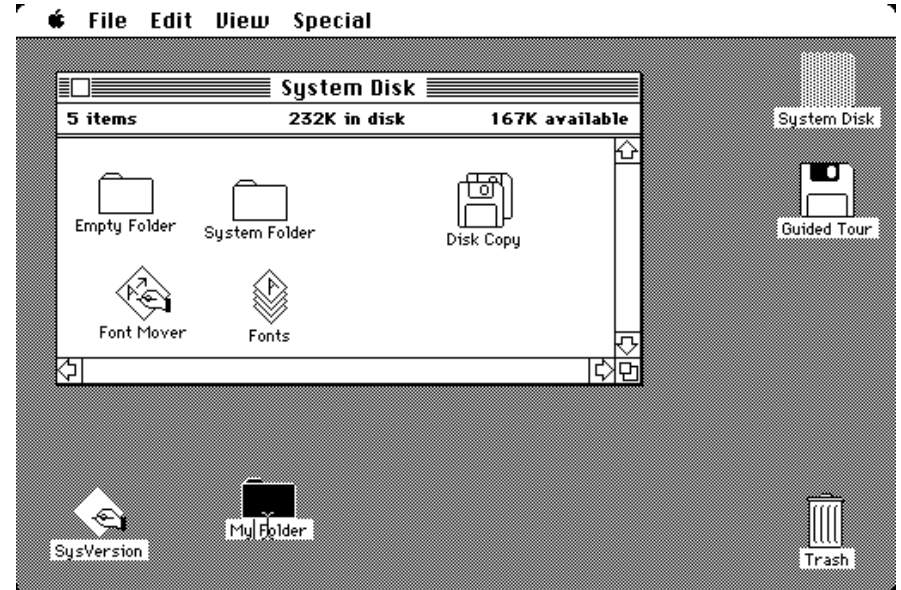
Timesharing and multitasking

- ▶ To ensure fairness, use a *timesharing* scheme in which the CPU cycles through all jobs
 - ▶ Each job is given a fixed amount of CPU time
 - ▶ Switching from one running job (or process) to another is called a *context switch*
 - ▶ A process may relinquish its time if blocked on an I/O request



History of operating systems (iv)

► Text CRTs (1970s) to an early Mac (1984)



History of operating systems (v)

► Personal computer revolution (1970s/80s)



History of operating systems (vi)


- ▶ The war begins...



sell your stock....


History of operating systems (vii)

- ▶ World Wide Web and Internet revolution
(1990s/2000s)



History of operating systems (viii)

- ▶ Mobile revolution (2010s)




We interrupt this program...

- ▶ Software instructions are executed by a Central Processing Unit (CPU)
 - ▶ An external hardware event triggers an interrupt by signaling the CPU
 - ▶ e.g. mouse movement, keyboard event
 - ▶ Software triggers an interrupt by executing a system call
 - ▶ e.g. disk read, graphics output, printer output

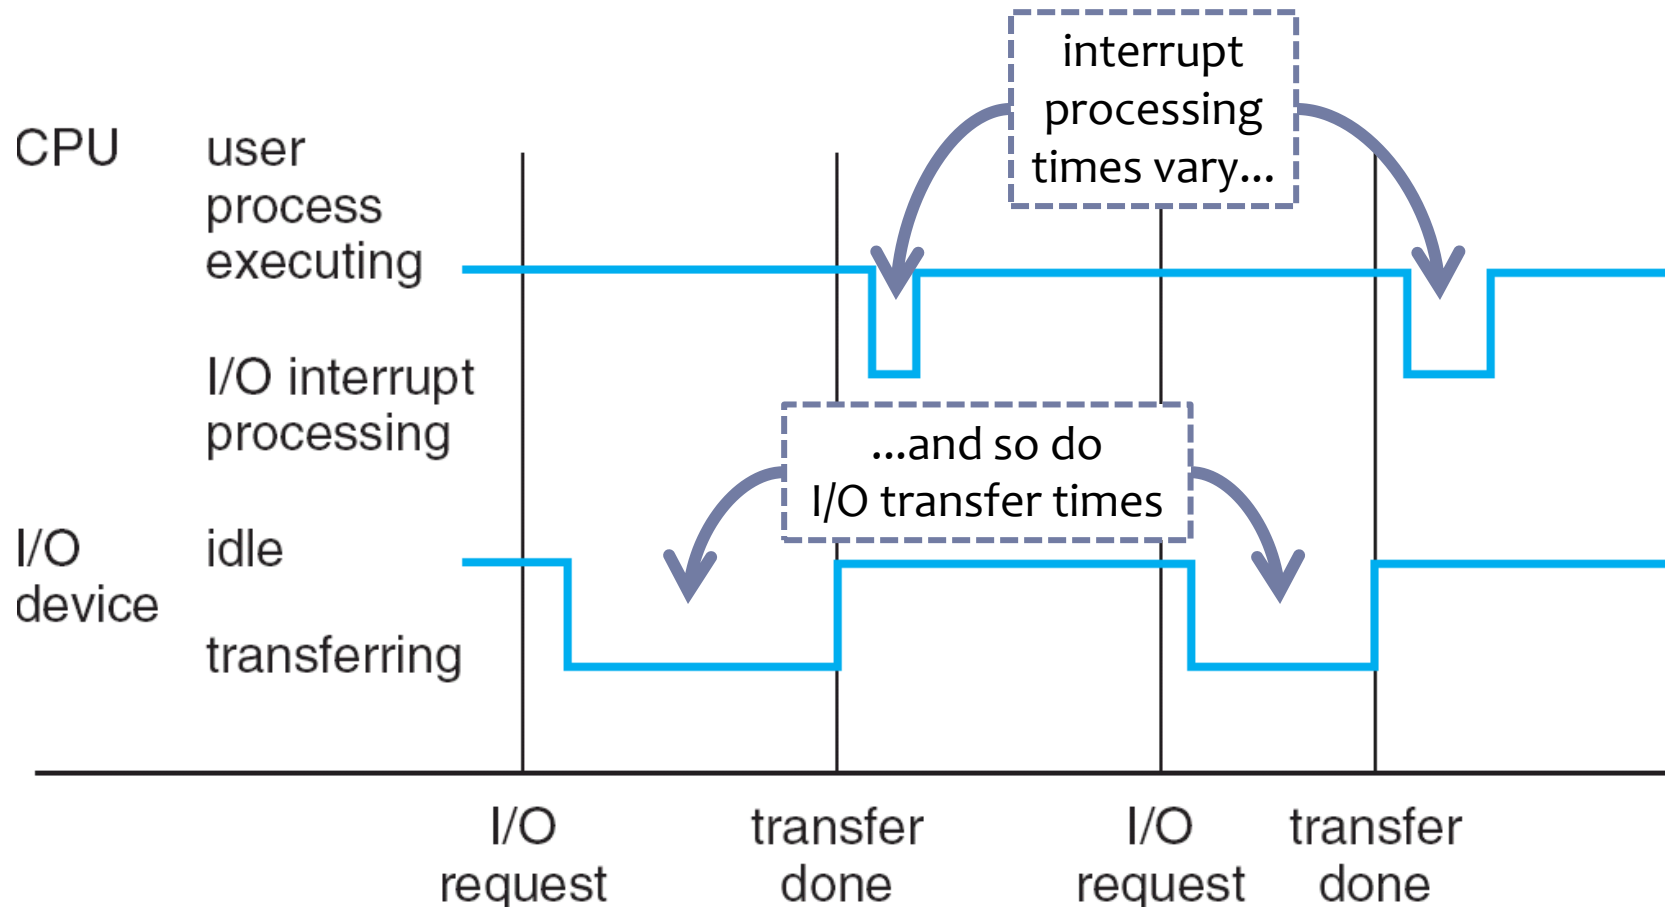


Interrupt mechanism

- ▶ Interrupts are handled much like calling a function in a programming language

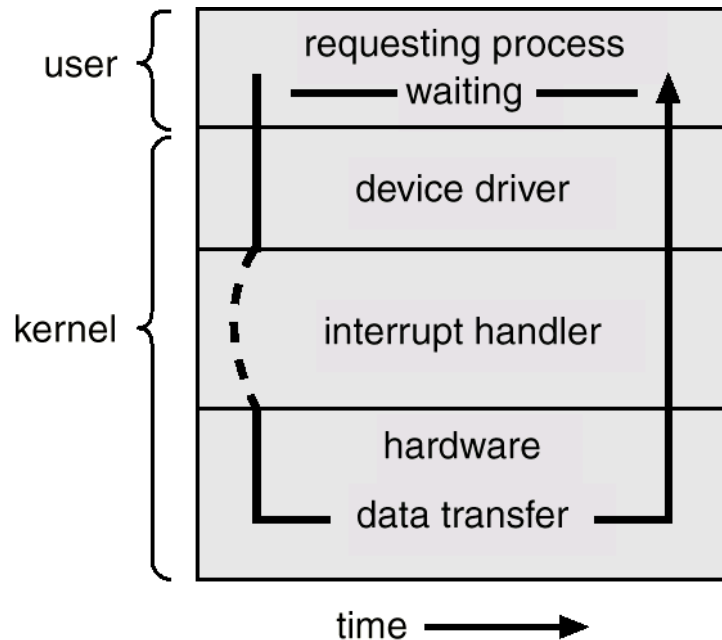


Typical interrupt timeline

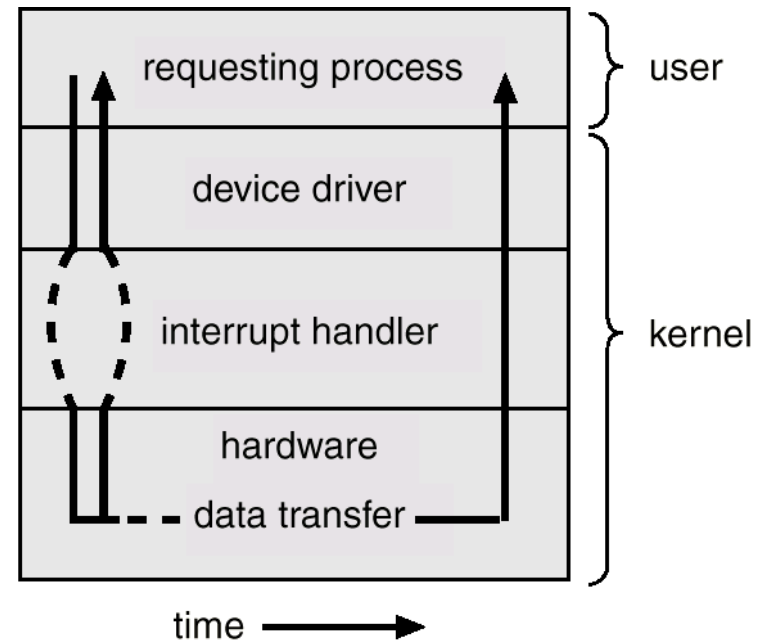


Synchronous and Asynchronous I/O

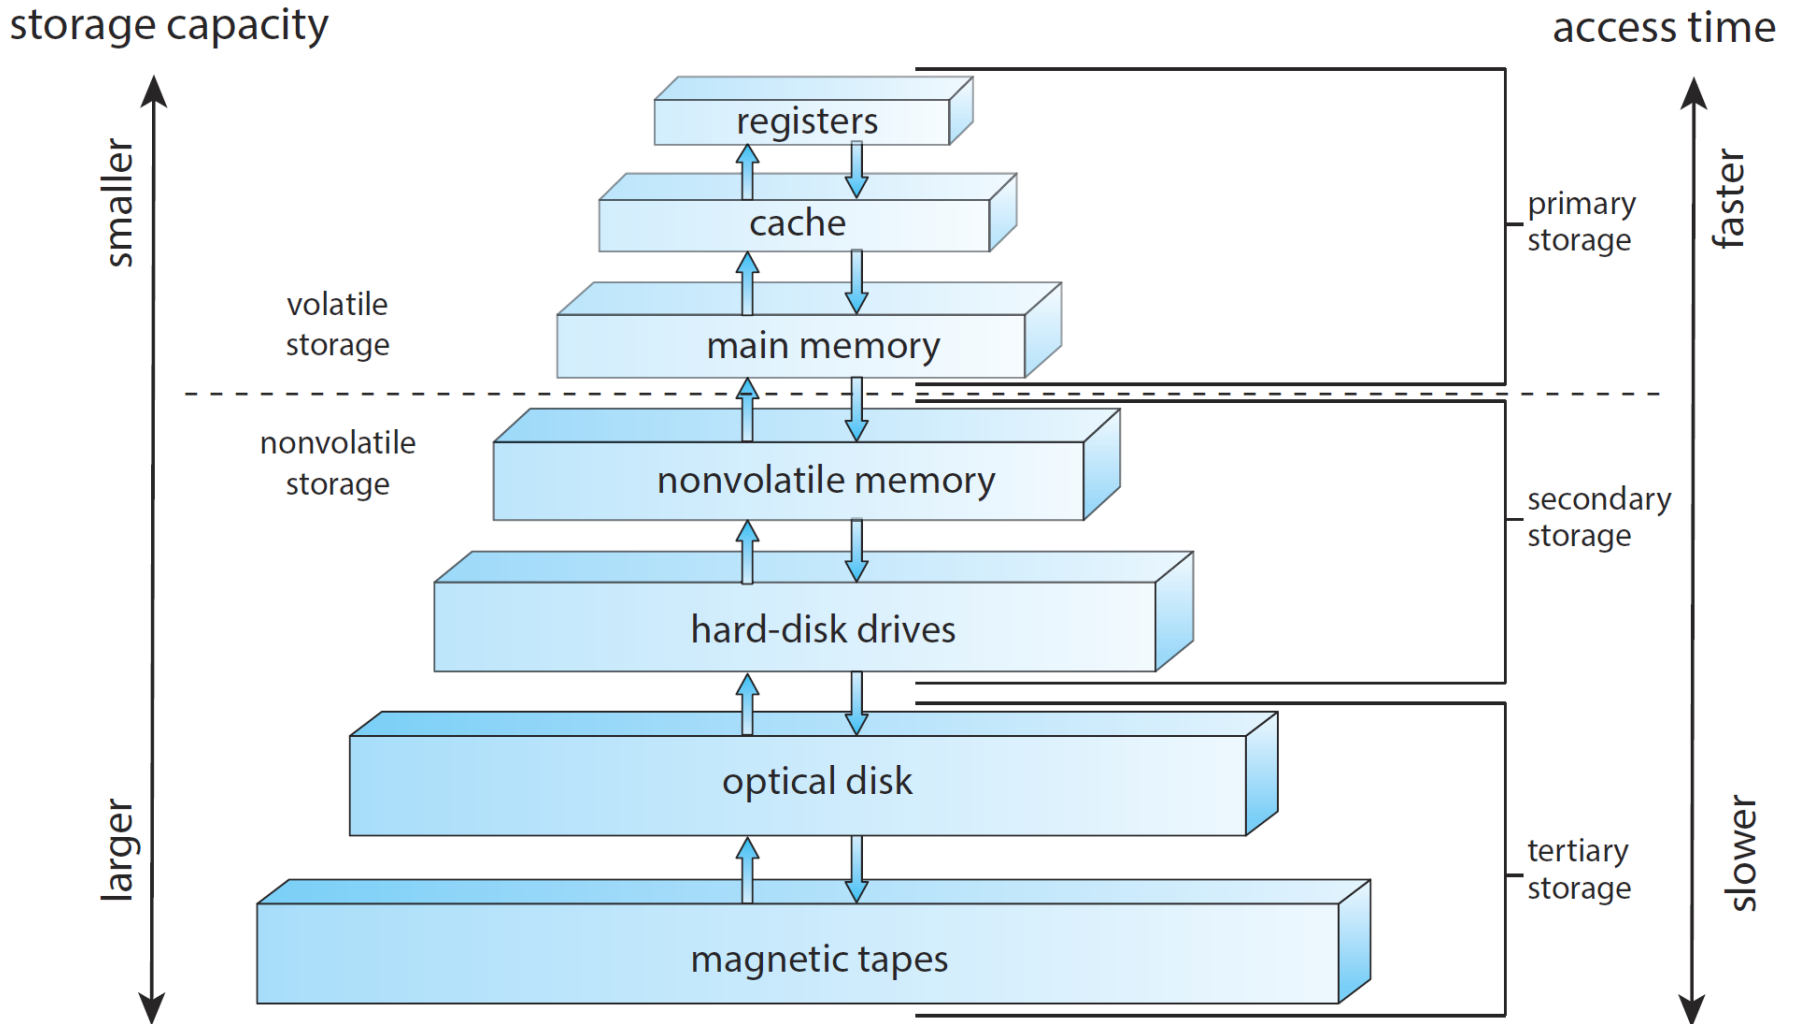
Synchronous



Asynchronous




Hierarchical storage architecture



Caching (i)

- ▶ Caching is a technique in which data is temporarily stored in a smaller and faster memory component
- ▶ Why implement caching in an operating system



Caching (ii)


- ▶ A key goal in operating system design is achieving *fast and efficient performance*

Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25 – 0.5	0.5 – 25	80 – 250	5,000.000
Bandwidth (MB/sec)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape



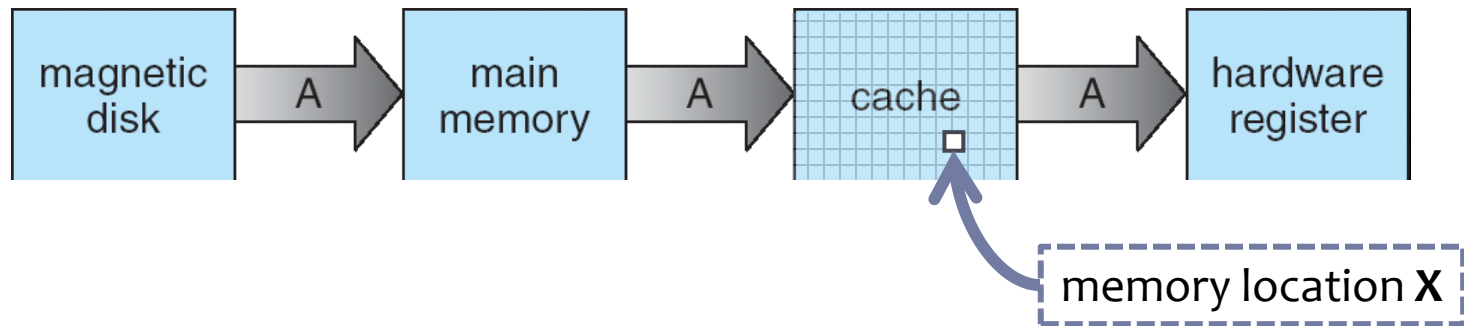
Caching (iii)

- ▶ What's the caching algorithm?
 - ▶ When the operating system attempts to read from memory, check to see if the requested data is already in the cache
 - ▶ If it is, data is read from the cache (*fast!*)
 - ▶ If not, data is copied from memory to the cache (*maybe next time...*)



Principle of locality

- ▶ When a running program reads from memory location **X**, the *principle of locality* predicts that the next memory location requested will be near **X**



- ▶ Store *pages* of data in a cache, where each page is typically the same size (e.g. 64KB)