

# Processes

Achmad Imam Kistijantoro ([imam@staff.stei.itb.ac.id](mailto:imam@staff.stei.itb.ac.id))

Judhi Santoso ([judhi@staff.stei.itb.ac.id](mailto:judhi@staff.stei.itb.ac.id))

Robithoh Annur ([robithoh@staff.stei.itb.ac.id](mailto:robithoh@staff.stei.itb.ac.id))

# Chapter 3: Processes

---

- ▶ Process Concept
- ▶ Process Scheduling
- ▶ Operations on Processes
- ▶ Cooperating Processes
- ▶ Interprocess Communication
- ▶ Communication in Client-Server Systems



# Process Concept

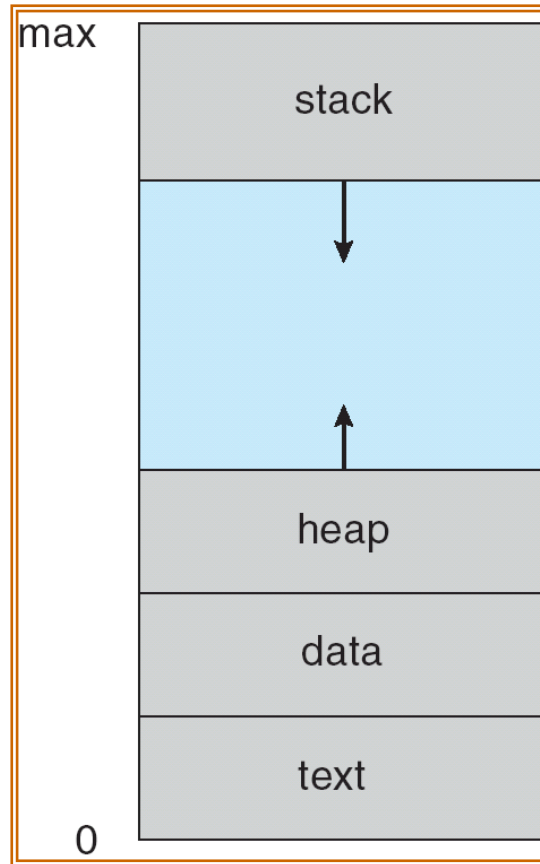
---

- ▶ An operating system executes a variety of programs:
  - ▶ Batch system – jobs
  - ▶ Time-shared systems – user programs or tasks
- ▶ Textbook uses the terms *job* and *process* almost interchangeably
- ▶ Process – a program in execution; process execution must progress in sequential fashion
- ▶ A process includes:
  - ▶ program counter
  - ▶ stack
  - ▶ data section



# Process in Memory

---



# Process State

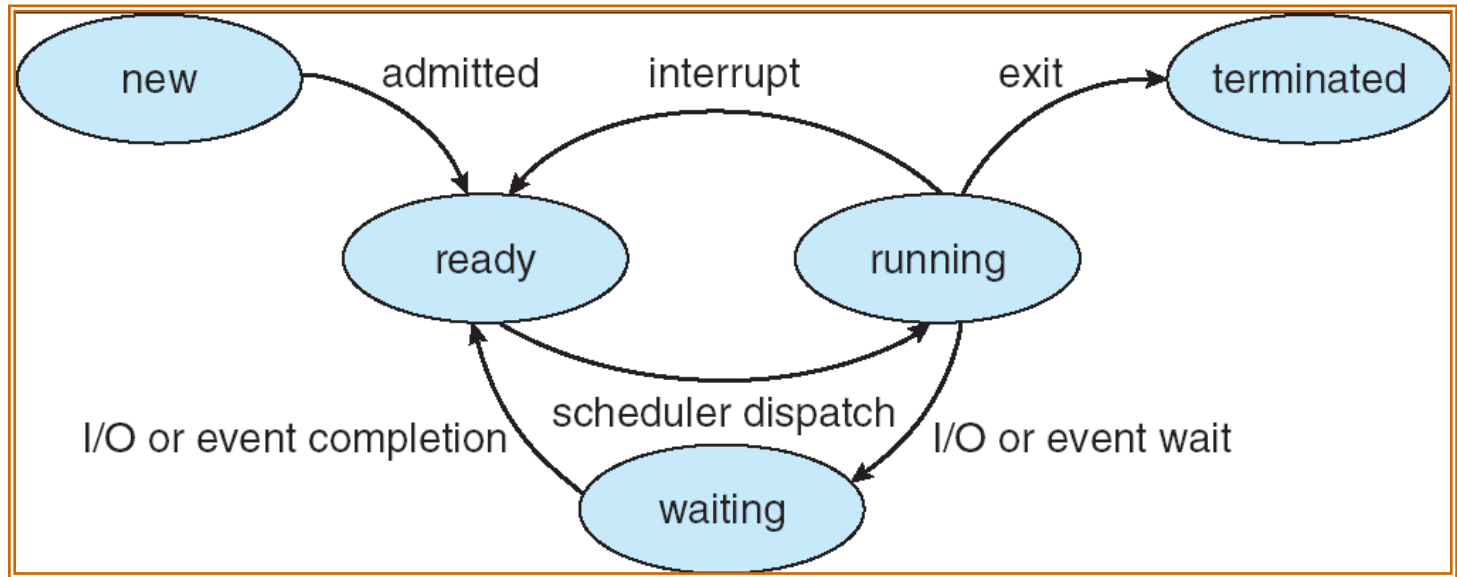
---

- ▶ As a process executes, it changes *state*
  - ▶ **new**: The process is being created
  - ▶ **running**: Instructions are being executed
  - ▶ **waiting**: The process is waiting for some event to occur
  - ▶ **ready**: The process is waiting to be assigned to a process
  - ▶ **terminated**: The process has finished execution



# Diagram of Process State

---



# Process Control Block (PCB)

---

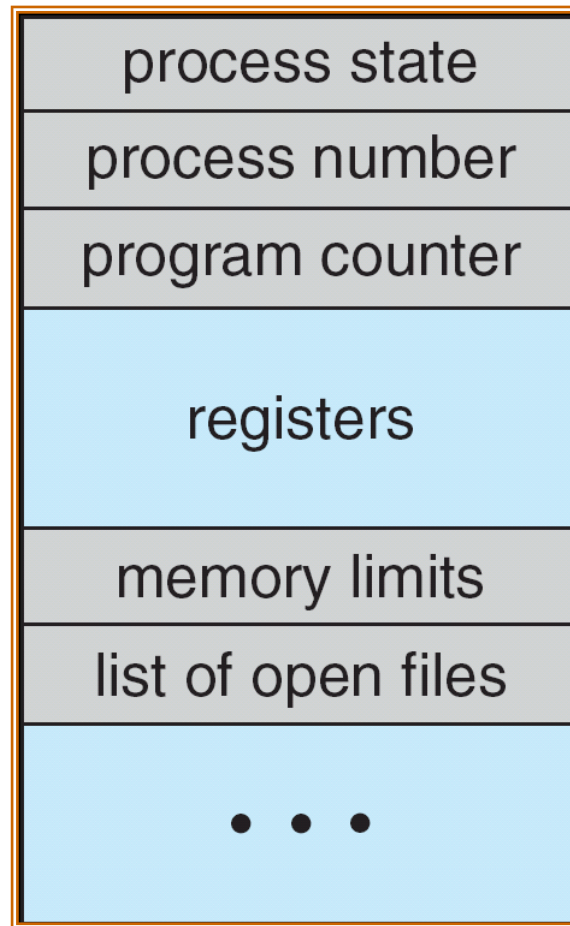
Information associated with each process

- ▶ Process state
- ▶ Program counter
- ▶ CPU registers
- ▶ CPU scheduling information
- ▶ Memory-management information
- ▶ Accounting information
- ▶ I/O status information



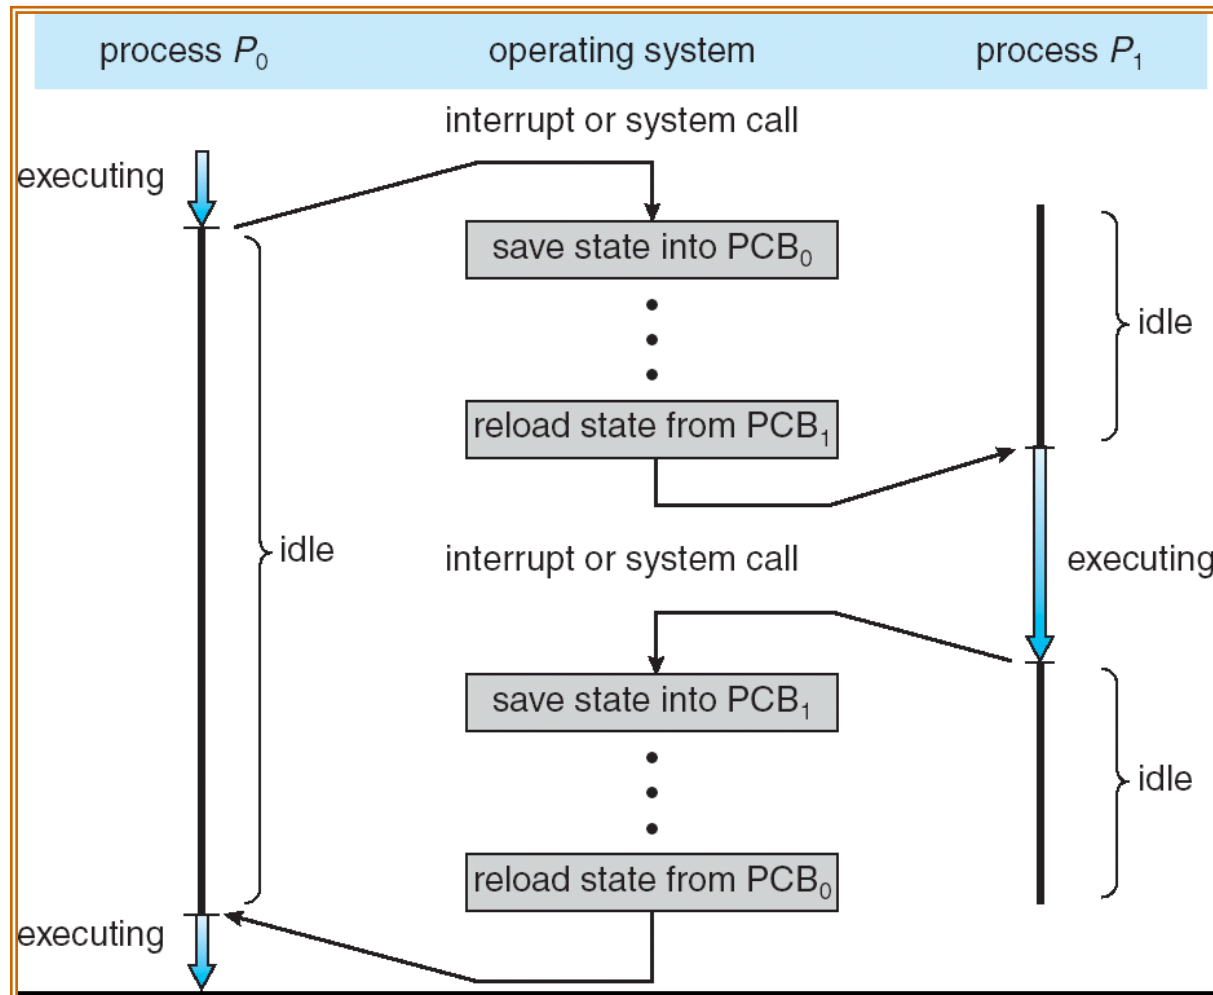
# Process Control Block (PCB)

---





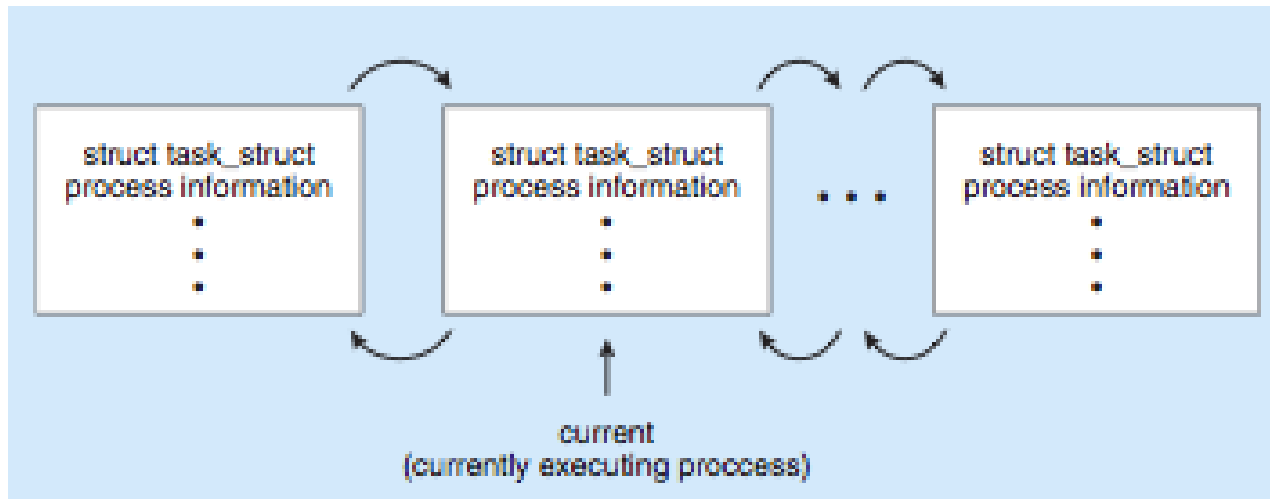
# CPU Switch From Process to Process



# Process Representation in Linux

- Represented by the C structure `task_struct`

```
pid t_pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



<https://git.kernel.org/pub/scm/linux/kernel/git/torvalds/linux.git/tree/include/linux/sched.h?h=v5.11-rc5#n649>

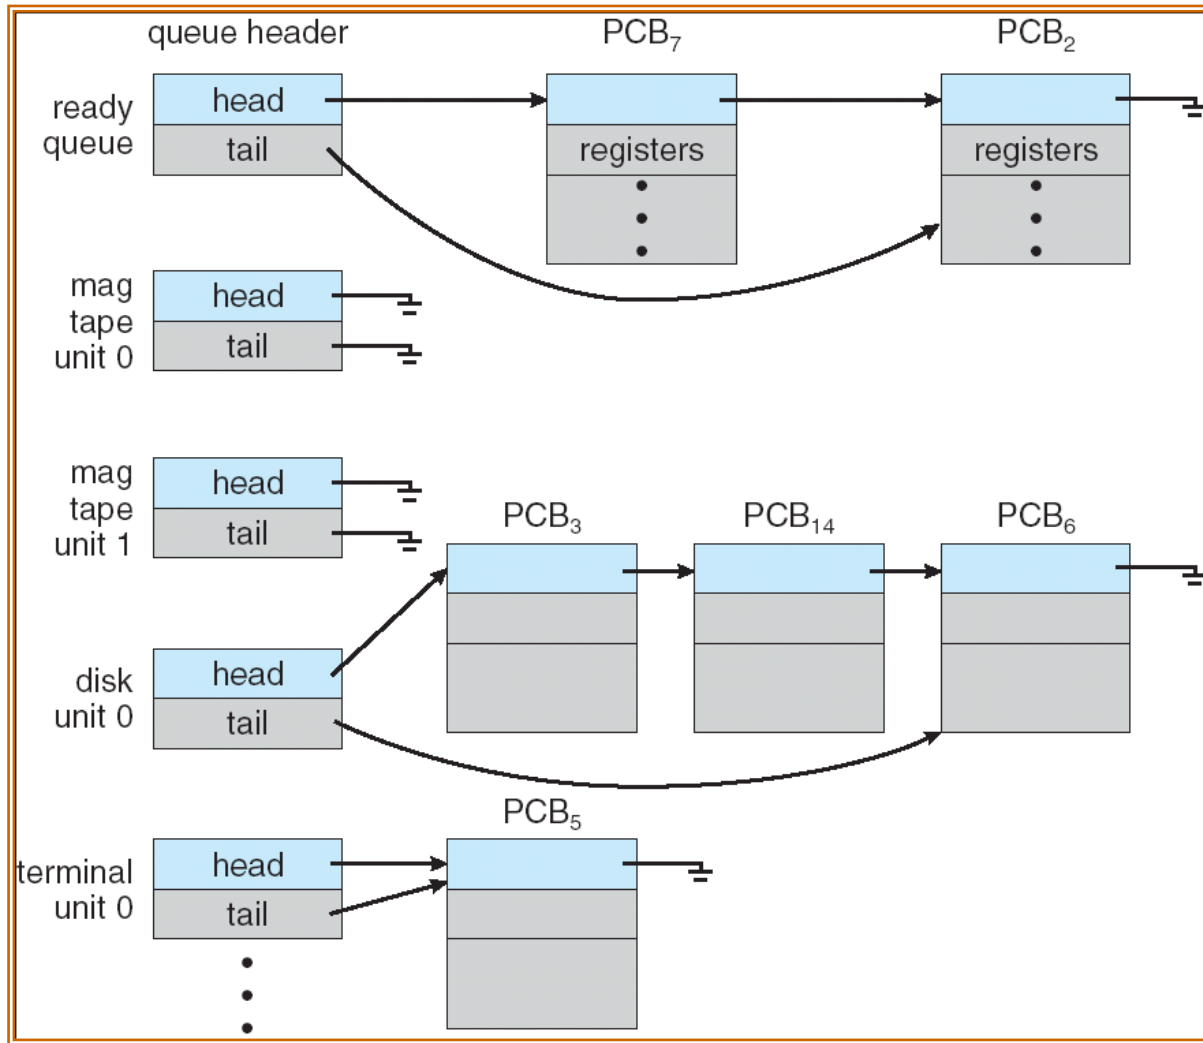
# Process Scheduling Queues

---

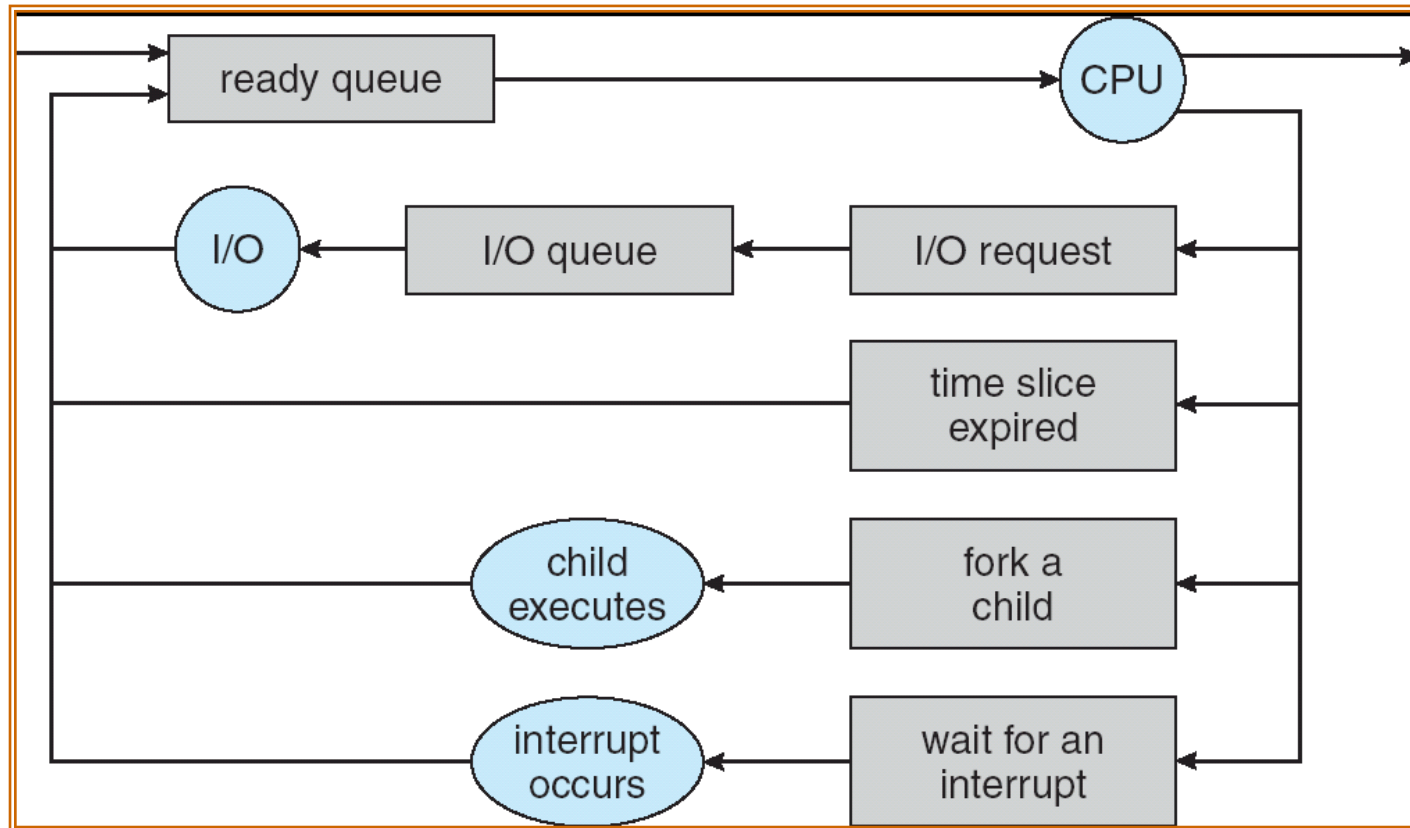
- ▶ **Job queue** – set of all processes in the system
- ▶ **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- ▶ **Device queues** – set of processes waiting for an I/O device
- ▶ Processes migrate among the various queues



# Ready Queue And Various I/O Device Queues



# Representation of Process Scheduling



# Schedulers


---

- ▶ **Long-term scheduler** (or job scheduler)
  - selects which processes should be brought into the ready queue
- ▶ **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU



# Addition of Medium Term Scheduling

---



## Schedulers (Cont.)

---

- ▶ Short-term scheduler is invoked very frequently (milliseconds)  $\Rightarrow$  (must be fast)
- ▶ Long-term scheduler is invoked very infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)
- ▶ The long-term scheduler controls the *degree of multiprogramming*
- ▶ Processes can be described as either:
  - ▶ **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
  - ▶ **CPU-bound process** – spends more time doing computations; few very long CPU bursts





# Multitasking in Mobile Systems

---

- ▶ Some systems / early systems allow only one process to run, others suspended
- ▶ Due to screen real estate, user interface limits iOS provides for a
  - ▶ Single **foreground** process- controlled via user interface
  - ▶ Multiple **background** processes– in memory, running, but not on the display, and with limits
  - ▶ Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- ▶ Android runs foreground and background, with fewer limits
  - ▶ Background process uses a **service** to perform tasks
  - ▶ Service can keep running even if background process is suspended
  - ▶ Service has no user interface, small memory use



# Context Switch

---

- ▶ When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process
- ▶ Context-switch time is overhead; the system does no useful work while switching
- ▶ Time dependent on hardware support



# Process Creation

---

- ▶ Parent process create children processes, which, in turn create other processes, forming a tree of processes
- ▶ Resource sharing
  - ▶ Parent and children share all resources
  - ▶ Children share subset of parent's resources
  - ▶ Parent and child share no resources
- ▶ Execution
  - ▶ Parent and children execute concurrently
  - ▶ Parent waits until children terminate



# Process Creation (Cont.)


---

- ▶ **Address space**
  - ▶ Child duplicate of parent
  - ▶ Child has a program loaded into it
- ▶ **UNIX examples**
  - ▶ **fork** system call creates new process
  - ▶ **exec** system call used after a **fork** to replace the process' memory space with a new program



# Process Creation

---



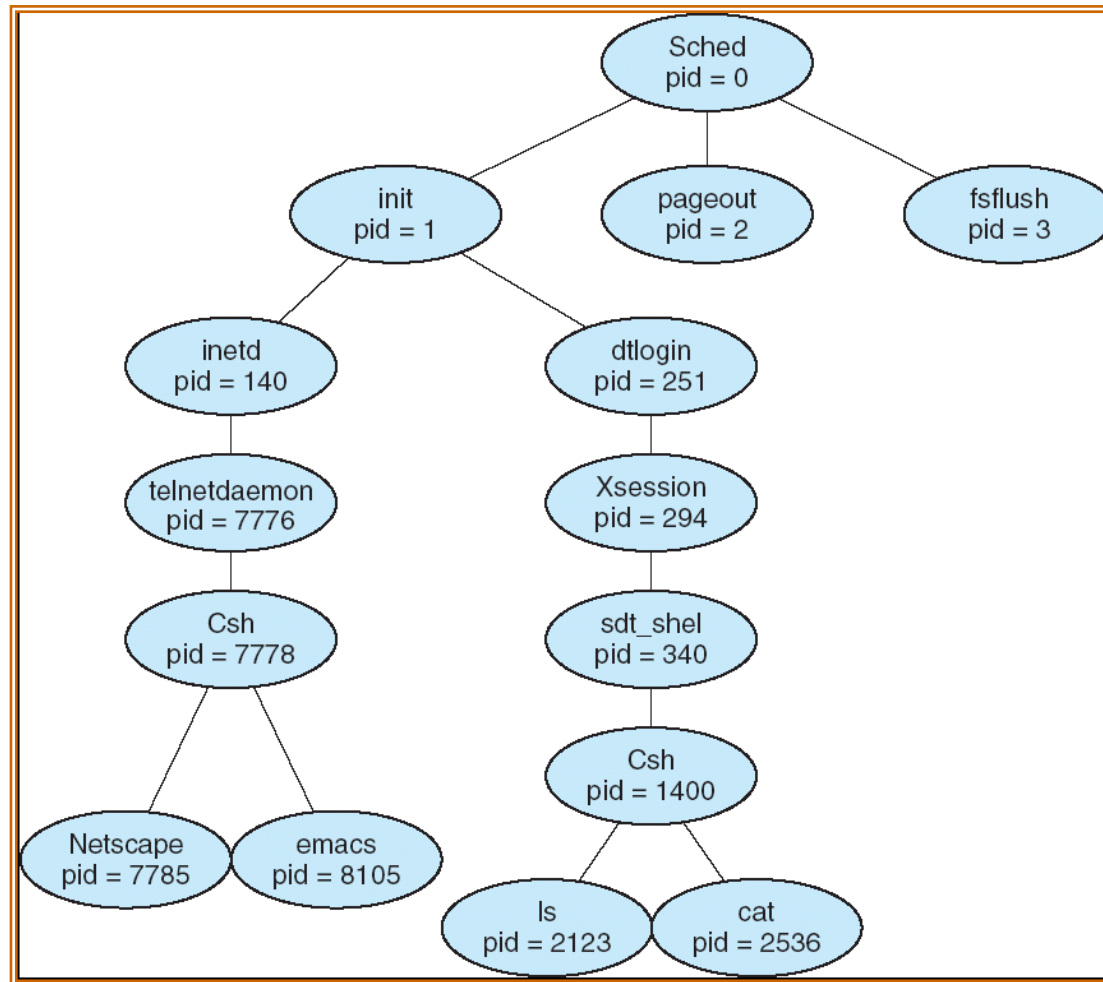
# C Program Forking Separate Process

---

```
int main()
{
    Pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait (NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```



# A tree of processes on a typical Solaris



# Process Termination

---

- ▶ Process executes last statement and asks the operating system to delete it (**exit**)
  - ▶ Output data from child to parent (via **wait**)
  - ▶ Process' resources are deallocated by operating system
- ▶ Parent may terminate execution of children processes (**abort**)
  - ▶ Child has exceeded allocated resources
  - ▶ Task assigned to child is no longer required
  - ▶ If parent is exiting
    - ▶ Some operating system do not allow child to continue if its parent terminates
      - All children terminated - *cascading termination*





# Cooperating Processes

---

- ▶ **Independent** process cannot affect or be affected by the execution of another process
- ▶ **Cooperating** process can affect or be affected by the execution of another process
- ▶ Advantages of process cooperation
  - ▶ Information sharing
  - ▶ Computation speed-up
  - ▶ Modularity
  - ▶ Convenience



# Producer-Consumer Problem

---

- ▶ Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
  - ▶ *unbounded-buffer* places no practical limit on the size of the buffer
  - ▶ *bounded-buffer* assumes that there is a fixed buffer size



# Bounded-Buffer – Shared-Memory Solution

---

- ▶ Shared data

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
    ...
```

```
} item;
```

```
item buffer[BUFFER_SIZE];
```

```
int in = 0;
```

```
int out = 0;
```

- ▶ Solution is correct, but can only use `BUFFER_SIZE-1` elements



# Bounded-Buffer – Insert() Method

---

```
while (true) {  
    /* Produce an item */  
    while (((in = (in + 1) % BUFFER SIZE) == out)  
        ; /* do nothing -- no free buffers */  
    buffer[in] = item;  
    in = (in + 1) % BUFFER SIZE;  
}
```



# Bounded Buffer – Remove() Method

---

```
while (true) {  
    while (in == out)  
        ; // do nothing -- nothing to  
        consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER SIZE;  
    return item;  
}
```



# Interprocess Communication (IPC)

---

- ▶ Mechanism for processes to communicate and to synchronize their actions
- ▶ Message system – processes communicate with each other without resorting to shared variables
- ▶ IPC facility provides two operations:
  - ▶ **send**(*message*) – message size fixed or variable
  - ▶ **receive**(*message*)
- ▶ If *P* and *Q* wish to communicate, they need to:
  - ▶ establish a *communication link* between them
  - ▶ exchange messages via send/receive
- ▶ Implementation of communication link
  - ▶ physical (e.g., shared memory, hardware bus)
  - ▶ logical (e.g., logical properties)



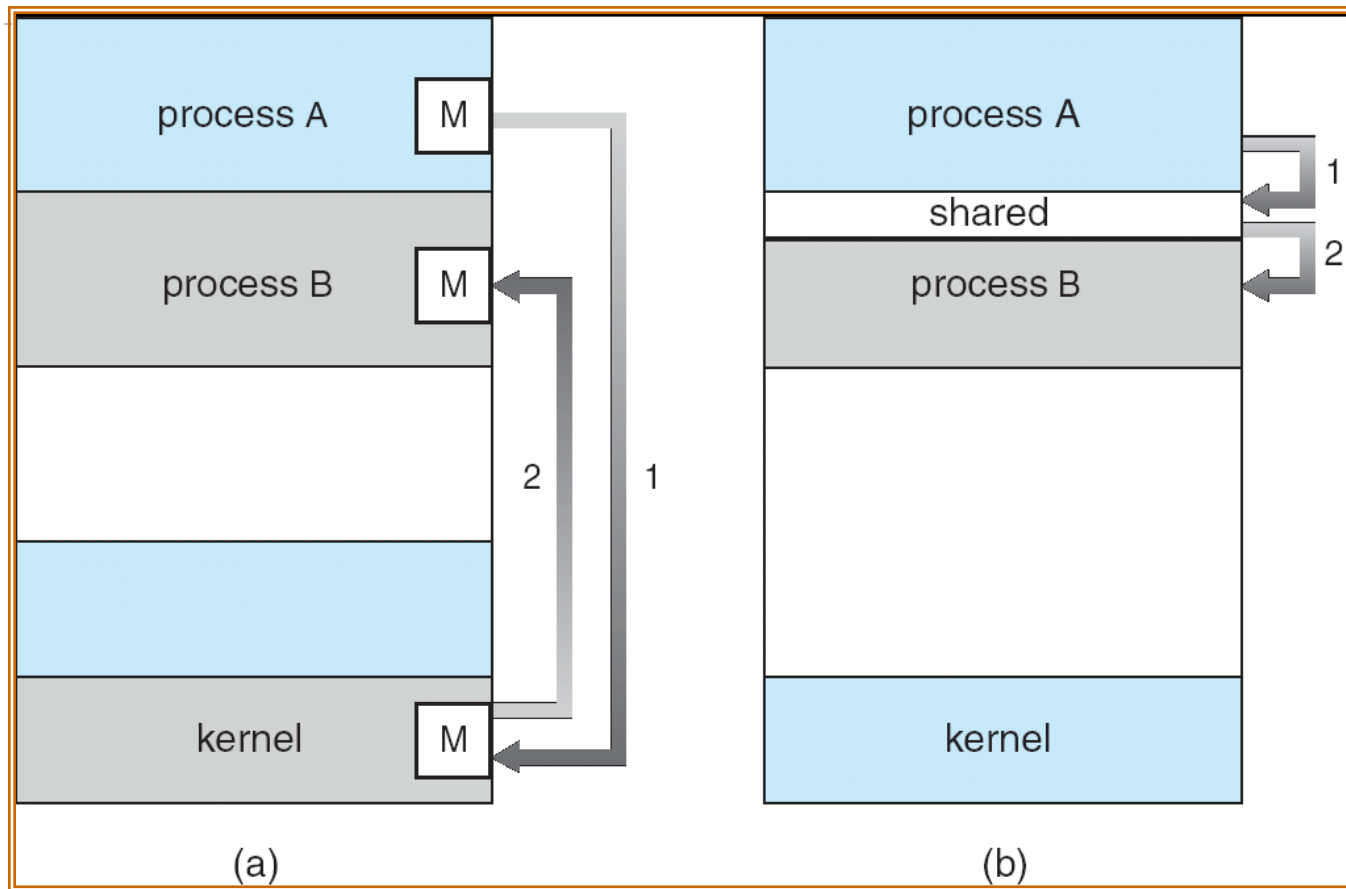
# Implementation Questions

---

- ▶ How are links established?
- ▶ Can a link be associated with more than two processes?
- ▶ How many links can there be between every pair of communicating processes?
- ▶ What is the capacity of a link?
- ▶ Is the size of a message that the link can accommodate fixed or variable?
- ▶ Is a link unidirectional or bi-directional?



# Communications Models





# Direct Communication

---

- ▶ Processes must name each other explicitly:
  - ▶ **send** ( $P, message$ ) – send a message to process  $P$
  - ▶ **receive**( $Q, message$ ) – receive a message from process  $Q$
- ▶ Properties of communication link
  - ▶ Links are established automatically
  - ▶ A link is associated with exactly one pair of communicating processes
  - ▶ Between each pair there exists exactly one link
  - ▶ The link may be unidirectional, but is usually bi-directional



# Indirect Communication

---

- ▶ Messages are directed and received from mailboxes (also referred to as ports)
  - ▶ Each mailbox has a unique id
  - ▶ Processes can communicate only if they share a mailbox
- ▶ Properties of communication link
  - ▶ Link established only if processes share a common mailbox
  - ▶ A link may be associated with many processes
  - ▶ Each pair of processes may share several communication links
  - ▶ Link may be unidirectional or bi-directional



# Indirect Communication

---

- ▶ Operations

- ▶ create a new mailbox
- ▶ send and receive messages through mailbox
- ▶ destroy a mailbox

- ▶ Primitives are defined as:

**send**( $A, message$ ) – send a message to mailbox  $A$

**receive**( $A, message$ ) – receive a message from mailbox  $A$



# Indirect Communication

---

## ▶ Mailbox sharing

- ▶  $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
- ▶  $P_1$  sends;  $P_2$  and  $P_3$  receive
- ▶ Who gets the message?

## ▶ Solutions

- ▶ Allow a link to be associated with at most two processes
- ▶ Allow only one process at a time to execute a receive operation
- ▶ Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.



# Synchronization

---

- ▶ Message passing may be either blocking or non-blocking
- ▶ **Blocking** is considered **synchronous**
  - ▶ **Blocking send** has the sender block until the message is received
  - ▶ **Blocking receive** has the receiver block until a message is available
- ▶ **Non-blocking** is considered **asynchronous**
  - ▶ **Non-blocking send** has the sender send the message and continue
  - ▶ **Non-blocking receive** has the receiver receive a valid message or null



# Buffering

---

- ▶ Queue of messages attached to the link; implemented in one of three ways
  1. Zero capacity – 0 messages  
Sender must wait for receiver (rendezvous)
  2. Bounded capacity – finite length of  $n$  messages  
Sender must wait if link full
  3. Unbounded capacity – infinite length  
Sender never waits



# Client-Server Communication

---

- ▶ Sockets
- ▶ Remote Procedure Calls
- ▶ Remote Method Invocation (Java)



# Sockets

---

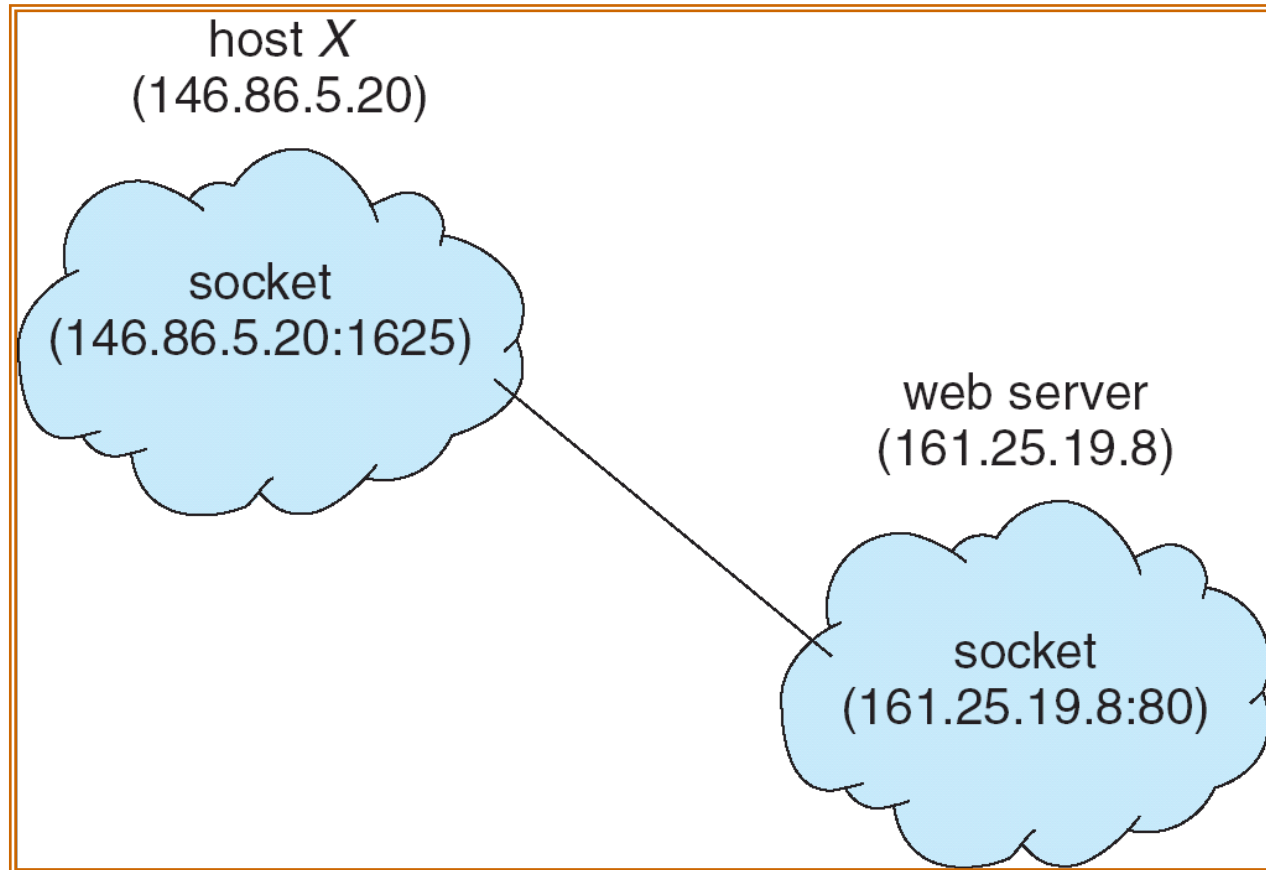
- ▶ A socket is defined as an *endpoint for communication*
- ▶ Concatenation of IP address and port
- ▶ The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- ▶ Communication consists between a pair of sockets





# Socket Communication

---



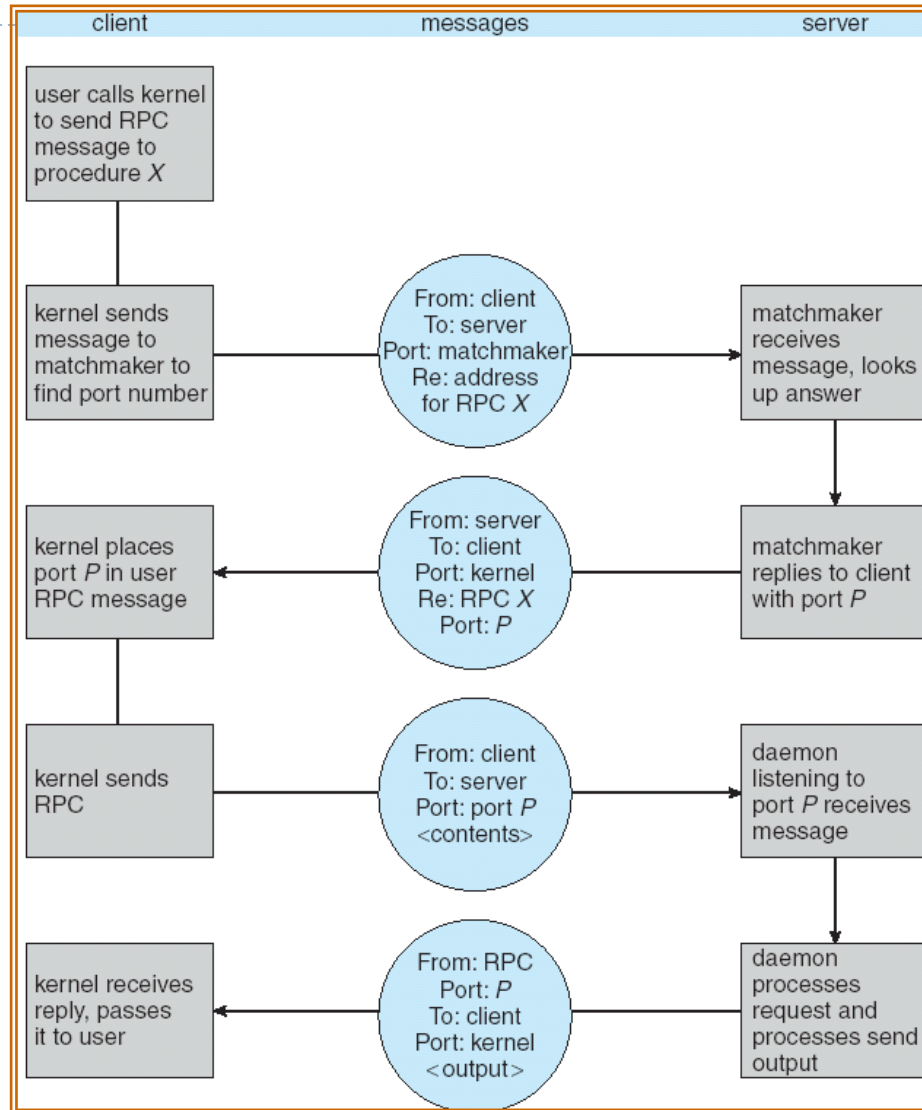
# Remote Procedure Calls

---

- ▶ Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- ▶ **Stubs** – client-side proxy for the actual procedure on the server.
- ▶ The client-side stub locates the server and *marshalls* the parameters.
- ▶ The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server.




# Execution of RPC




# Remote Method Invocation

---

- ▶ Remote Method Invocation (RMI) is a Java mechanism similar to RPCs.
- ▶ RMI allows a Java program on one machine to invoke a method on a remote object.



# Marshalling Parameters



# Contoh IPC: shared memory di Linux

---

- ▶ Kernel mengelola struktur untuk segmen shared memory



```
struct shmid_ds {
    struct ipc_perm shm_perm;    /* see Section 15.6.2 */
    size_t          shm_segsz;   /* size of segment in bytes */
    pid_t           shm_lpid;    /* pid of last shmop() */
    pid_t           shm_cpid;    /* pid of creator */
    shmatt_t        shm_nattch;  /* number of current attaches */
    time_t          shm_atime;   /* last-attach time */
    time_t          shm_dtime;   /* last-detach time */
    time_t          shm_ctime;   /* last-change time */
    .
    .
    .
};
```



# Contoh IPC: shared memory di Linux

---

- ▶ `shmget`: untuk mendapatkan shared memory id
- ▶ Key dapat menggunakan `IPC_PRIVATE`, atau nilai yang disepakati oleh kedua proses, atau menggunakan `ftok(path, id)` untuk membangkitkan key dari path dan id yang disepakati

```
#include <sys/shm.h>

int shmget(key_t key, size_t size, int flag);
```

Returns: shared memory ID if OK, 1 on error



# Contoh IPC: shared memory di Linux

---

- ▶ **shmctl**: menyediakan berbagai control untuk shared memory
- ▶ **cmd**:
  - ▶ **IPC\_STAT**: membaca `shmid_ds`
  - ▶ **IPC\_SET**: menset nilai `shm_perm.uid`, `shm_perm.gid` dan `shm_perm.mode`
  - ▶ **IPC\_RMID**: menghapus shared memory segment
  - ▶ **SHM\_LOCK** dan **SHM\_UNLOCK**

```
#include <sys/shm.h>

int shmctl(int shmid, int cmd, struct shmid_ds *buf);
```

Returns: 0 if OK, 1 on error



# Contoh IPC: shared memory di Linux

---

- ▶ `shmat`: attach shared memory ke process address space
- ▶ `Shmdt`: detach shared memory

```
#include <sys/shm.h>
```

```
void *shmat(int shmid, const void *addr, int flag);
```

Returns: pointer to shared memory segment if OK, 1 on error

```
#include <sys/shm.h>
```

```
int shmdt(void *addr);
```

Returns: 0 if OK, 1 on error

```

#include "apue.h"
#include <sys/shm.h>

#define ARRAY_SIZE 40000
#define MALLOC_SIZE 100000
----- #define SHM_SIZE 100000 -----
#define SHM_MODE 0600 /* user read/write */

char array[ARRAY_SIZE]; /* uninitialized data = bss */

int
main(void)
{
    int shmid;
    char *ptr, *shmptr;

    printf("array[] from %lx to %lx\n", (unsigned long)&array[0],
           (unsigned long)&array[ARRAY_SIZE]);
    printf("stack around %lx\n", (unsigned long)&shmid);

    if ((ptr = malloc(MALLOC_SIZE)) == NULL)
        err_sys("malloc error");
    printf("malloced from %lx to %lx\n", (unsigned long)ptr,
           (unsigned long)ptr+MALLOC_SIZE);

    if ((shmid = shmget(IPC_PRIVATE, SHM_SIZE, SHM_MODE)) < 0)
        err_sys("shmget error");
    if ((shmptr = shmat(shmid, 0, 0)) == (void *)-1)
        err_sys("shmat error");
    printf("shared memory attached from %lx to %lx\n",
           (unsigned long)shmptr, (unsigned long)shmptr+SHM_SIZE);

    if (shmctl(shmid, IPC_RMID, 0) < 0)
        err_sys("shmctl error");

    exit(0);
----- }

```

# IPC POSIX Producer

---

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```



# IPC POSIX Consumer

---

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE = 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

---



# Contoh IPC: message queue

---

## ► Struktur di kernel

```
struct msqid_ds {
    struct ipc_perm  msg_perm;           /* see Section 15.6.2 */
    msgqnum_t        msg_qnum;           /* # of messages on queue */
    msglen_t         msg_qbytes;         /* max # of bytes on queue */
    pid_t            msg_lspid;          /* pid of last msgsnd() */
    pid_t            msg_lrpid;          /* pid of last msgrcv() */
    time_t           msg_stime;          /* last-msgsnd() time */
    time_t           msg_rtime;          /* last-msgrcv() time */
    time_t           msg_ctime;          /* last-change time */
    .
    .
    .
};
```



- 
- ▶ **Msgget:** membangkitkan id dari key

```
#include <sys/msg.h>

int msgget(key_t key, int flag);
```

Returns: message queue ID if OK, 1 on error



---

## ► Mengirim dan menerima data ke/dari message queue

[\[View full width\]](#)

```
#include <sys/msg.h>

int msgsnd(int msqid, const void *ptr, size_t
↳ nbytes, int flag);
```

Returns: 0 if OK, 1 on error

[\[View full width\]](#)

```
#include <sys/msg.h>

ssize_t msgrcv(int msqid, void *ptr, size_t nbytes
↳ , long type, int flag);
```

Returns: size of data portion of message if OK, 1 on error

---