


# IF2230

## I/O Systems



# Chapter 12: I/O Systems

---

- ▶ I/O Hardware
- ▶ Application I/O Interface
- ▶ Kernel I/O Subsystem
- ▶ Transforming I/O Requests to Hardware Operations
- ▶ Streams
- ▶ Performance



# Objectives

---

- ▶ Explore the structure of an operating system's I/O subsystem
- ▶ Discuss the principles of I/O hardware and its complexity
- ▶ Provide details of the performance aspects of I/O hardware and software



# Overview

---

- ▶ I/O management is a major component of operating system design and operation
  - ▶ Important aspect of computer operation
  - ▶ I/O devices vary greatly
  - ▶ Various methods to control them
  - ▶ Performance management
  - ▶ New types of devices frequent
- ▶ Ports, busses, device controllers connect to various devices
- ▶ **Device drivers** encapsulate device details
  - ▶ Present uniform device-access interface to I/O subsystem



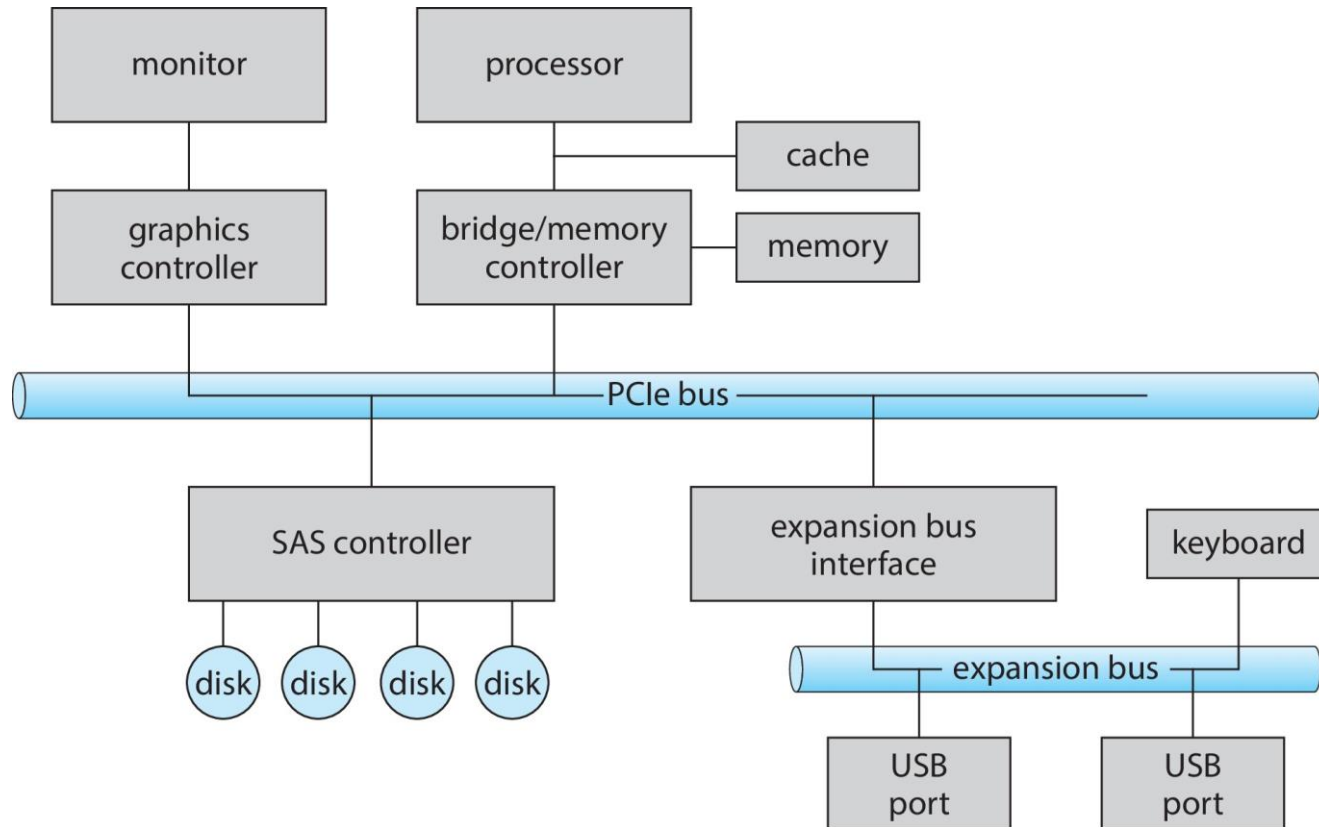
# I/O Hardware

---

- ▶ Incredible variety of I/O devices
- ▶ Common concepts
  - ▶ **Port**
  - ▶ **Bus** (daisy chain or shared direct access)
    - ▶ **PCI** bus common in PCs and servers, PCI Express (**PCIe**)
    - ▶ **expansion bus** connects relatively slow devices
    - ▶ **Serial-attached SCSI (SAS)** common disk interface
  - ▶ **Controller (host adapter)**
- ▶ I/O instructions control devices
- ▶ Devices have addresses, used by
  - ▶ Direct I/O instructions
  - ▶ **Memory-mapped I/O**



# A Typical PC Bus Structure



# Device I/O Port Locations on PCs (partial)

---

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)



# Polling

- 
- For each byte of I/O
    1. Read busy bit from status register until 0
    2. Host sets read or write bit and if write copies data into data-out register
    3. Host sets command-ready bit
    4. Controller sets busy bit, executes transfer
    5. Controller clears busy bit, error bit, command-ready bit when transfer done
  - Step 1 is **busy-wait** cycle to wait for I/O from device
    - Reasonable if device is fast
    - But inefficient if device slow
    - CPU switches to other tasks?
      - But if miss a cycle data overwritten / lost
- 



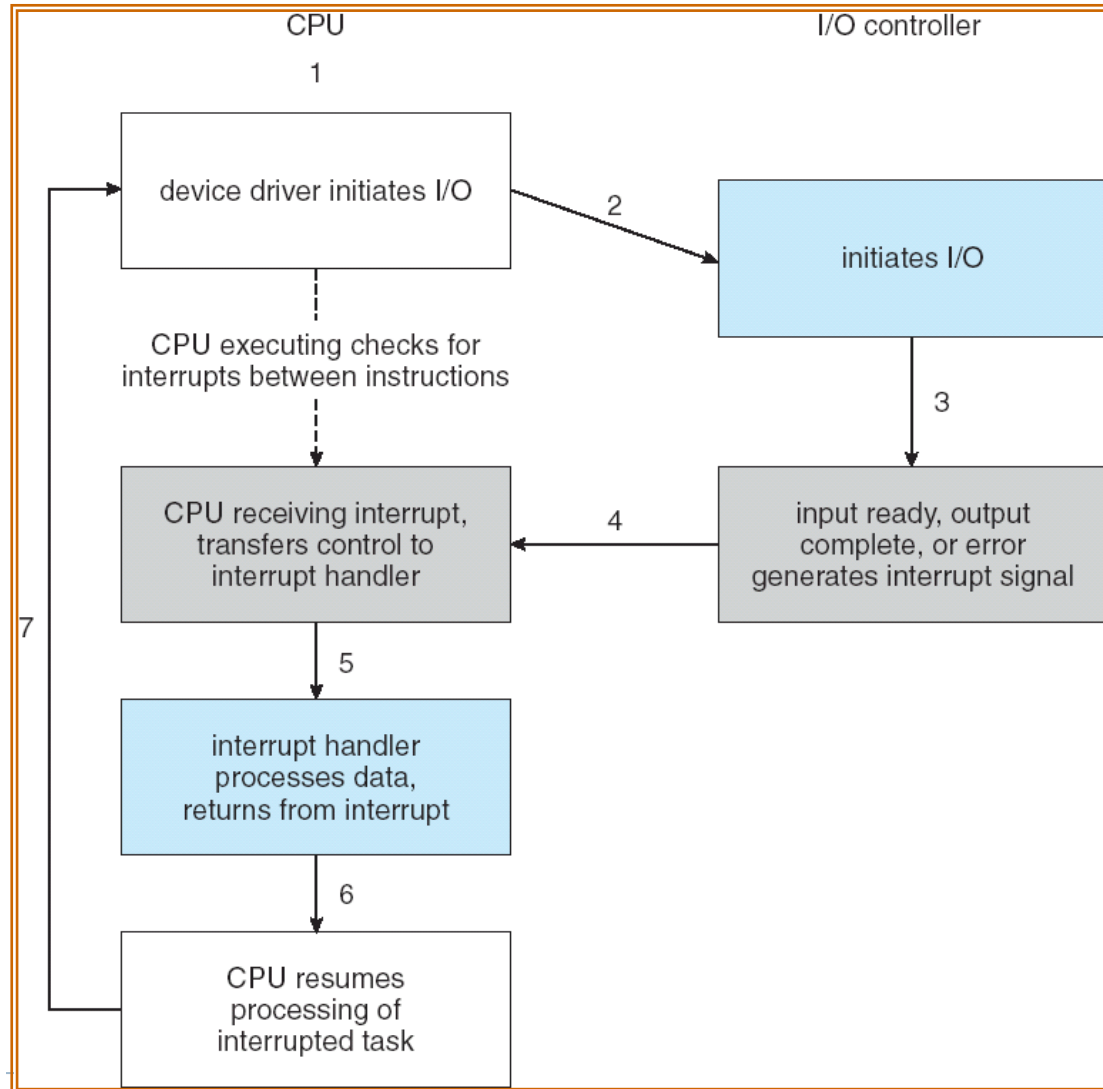


# Interrupts

- ▶ Polling can happen in 3 instruction cycles
  - ▶ Read status, logical-and to extract status bit, branch if not zero
  - ▶ How to be more efficient if non-zero infrequently?
- ▶ CPU **Interrupt-request line** triggered by I/O device
  - ▶ Checked by processor after each instruction
- ▶ **Interrupt handler** receives interrupts
  - ▶ **Maskable** to ignore or delay some interrupts
- ▶ **Interrupt vector** to dispatch interrupt to correct handler
  - ▶ Context switch at start and end
  - ▶ Based on priority
  - ▶ Some **nonmaskable**
  - ▶ Interrupt chaining if more than one device at same interrupt number



# Interrupt-Driven I/O Cycle



# Interrupts (Cont.)

---

- ▶ Interrupt mechanism also used for **exceptions**
  - ▶ Terminate process, crash system due to hardware error
- ▶ Page fault executes when memory access error
- ▶ System call executes via **trap** to trigger kernel to execute request
- ▶ Multi-CPU systems can process interrupts concurrently
  - ▶ If operating system designed to handle it
- ▶ Used for time-sensitive processing, frequent, must be fast



# Latency

- ▶ Stressing interrupt management because even single-user systems manage hundreds or interrupts per second and servers hundreds of thousands
- ▶ For example, a quiet macOS desktop generated 23,000 interrupts over 10 seconds

Fri Nov 25 13:55:59		0:00:10
	SCHEDULER	INTERRUPTS
-----	-----	-----
total_samples	13	22998
delays < 10 usecs	12	16243
delays < 20 usecs	1	5312
delays < 30 usecs	0	473
delays < 40 usecs	0	590
delays < 50 usecs	0	61
delays < 60 usecs	0	317
delays < 70 usecs	0	2
delays < 80 usecs	0	0
delays < 90 usecs	0	0
delays < 100 usecs	0	0
total < 100 usecs	13	22998

# Intel Pentium Processor Event-Vector Table

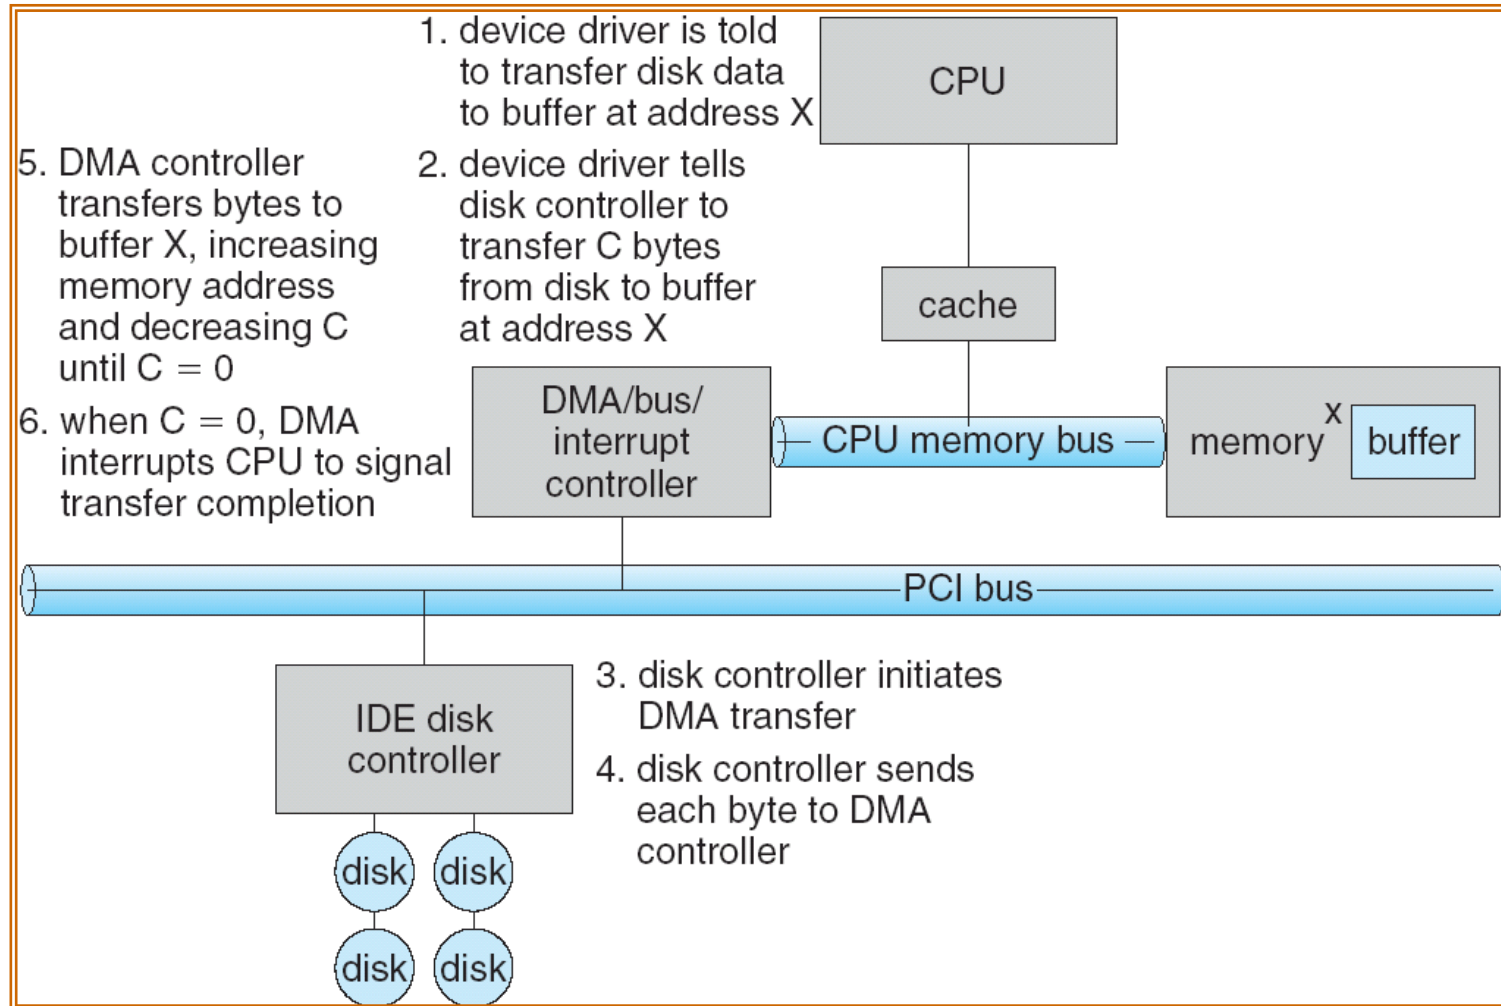
vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts

# Direct Memory Access

- ▶ Used to avoid **programmed I/O** (one byte at a time) for large data movement
- ▶ Requires **DMA** controller
- ▶ Bypasses CPU to transfer data directly between I/O device and memory
- ▶ OS writes DMA command block into memory
  - ▶ Source and destination addresses
  - ▶ Read or write mode
  - ▶ Count of bytes
  - ▶ Writes location of command block to DMA controller
  - ▶ Bus mastering of DMA controller – grabs bus from CPU
    - ▶ **Cycle stealing** from CPU but still much more efficient
  - ▶ When done, interrupts to signal completion
- ▶ Version that is aware of virtual addresses can be even more efficient - **DVMA**



# Six Step Process to Perform DMA Transfer



# Application I/O Interface

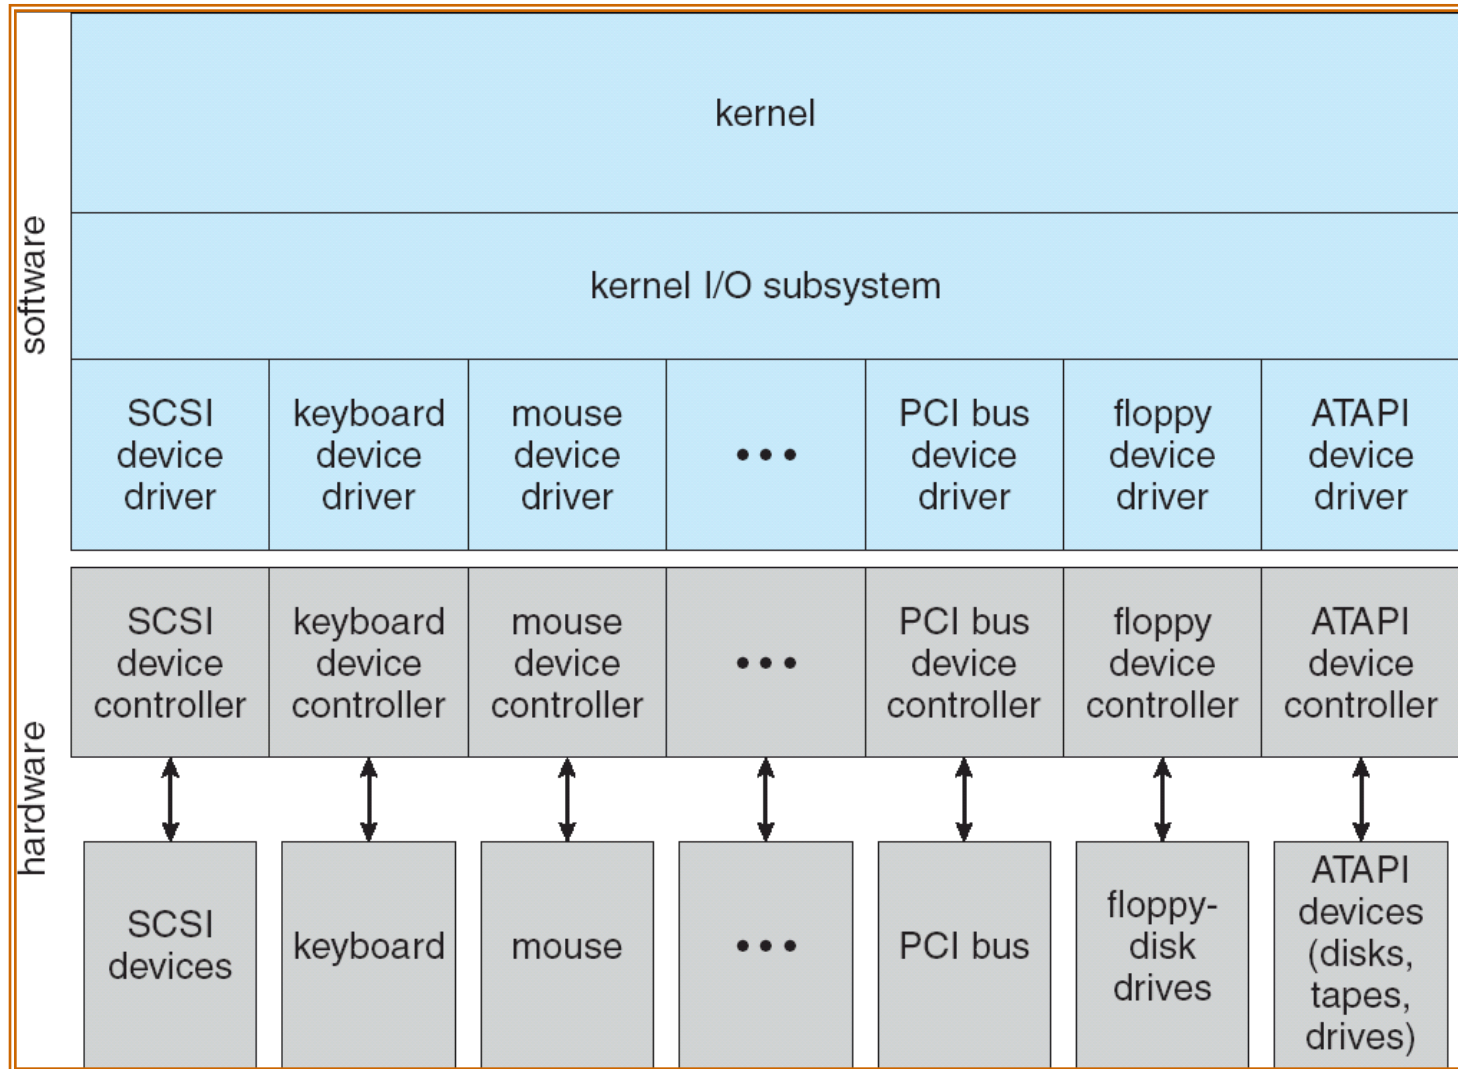
---

- ▶ I/O system calls encapsulate device behaviors in generic classes
- ▶ Device-driver layer hides differences among I/O controllers from kernel
- ▶ Devices vary in many dimensions
  - ▶ **Character-stream or block**
  - ▶ **Sequential or random-access**
  - ▶ **Sharable or dedicated**
  - ▶ **Speed of operation**
  - ▶ **read-write, read only, or write only**





# A Kernel I/O Structure



# Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read–write	CD-ROM graphics controller disk



# Characteristics of I/O Devices (Cont.)

- ▶ Subtleties of devices handled by device drivers
- ▶ Broadly I/O devices can be grouped by the OS into
  - ▶ Block I/O
  - ▶ Character I/O (Stream)
  - ▶ Memory-mapped file access
  - ▶ Network sockets
- ▶ For direct manipulation of I/O device specific characteristics, usually an escape / back door
  - ▶ Unix `ioctl()` call to send arbitrary bits to a device control register and data to device data register
- ▶ UNIX and Linux use tuple of “major” and “minor” device numbers to identify type and instance of devices (here major 8

```
brw-rw---- 1 root disk 8, 0 Mar 16 09:18 /dev/sda
brw-rw---- 1 root disk 8, 1 Mar 16 09:18 /dev/sda1
brw-rw---- 1 root disk 8, 2 Mar 16 09:18 /dev/sda2
brw-rw---- 1 root disk 8, 3 Mar 16 09:18 /dev/sda3
```

# Block and Character Devices

---

- ▶ Block devices include disk drives
  - ▶ Commands include read, write, seek
  - ▶ **Raw I/O**, **direct I/O**, or file-system access
  - ▶ Memory-mapped file access possible
    - ▶ File mapped to virtual memory and clusters brought via demand paging
  - ▶ DMA
- ▶ Character devices include keyboards, mice, serial ports
  - ▶ Commands include `get()` , `put()`
  - ▶ Libraries layered on top allow line editing



# Network Devices

---

- ▶ Varying enough from block and character to have own interface
- ▶ Linux, Unix, Windows and many others include **socket** interface
  - ▶ Separates network protocol from network operation
  - ▶ Includes `select()` functionality
- ▶ Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)



# Clocks and Timers

---

- ▶ Provide current time, elapsed time, timer
- ▶ **Programmable interval timer** used for timings, periodic interrupts
- ▶ `ioctl` (on UNIX) covers odd aspects of I/O such as clocks and timers



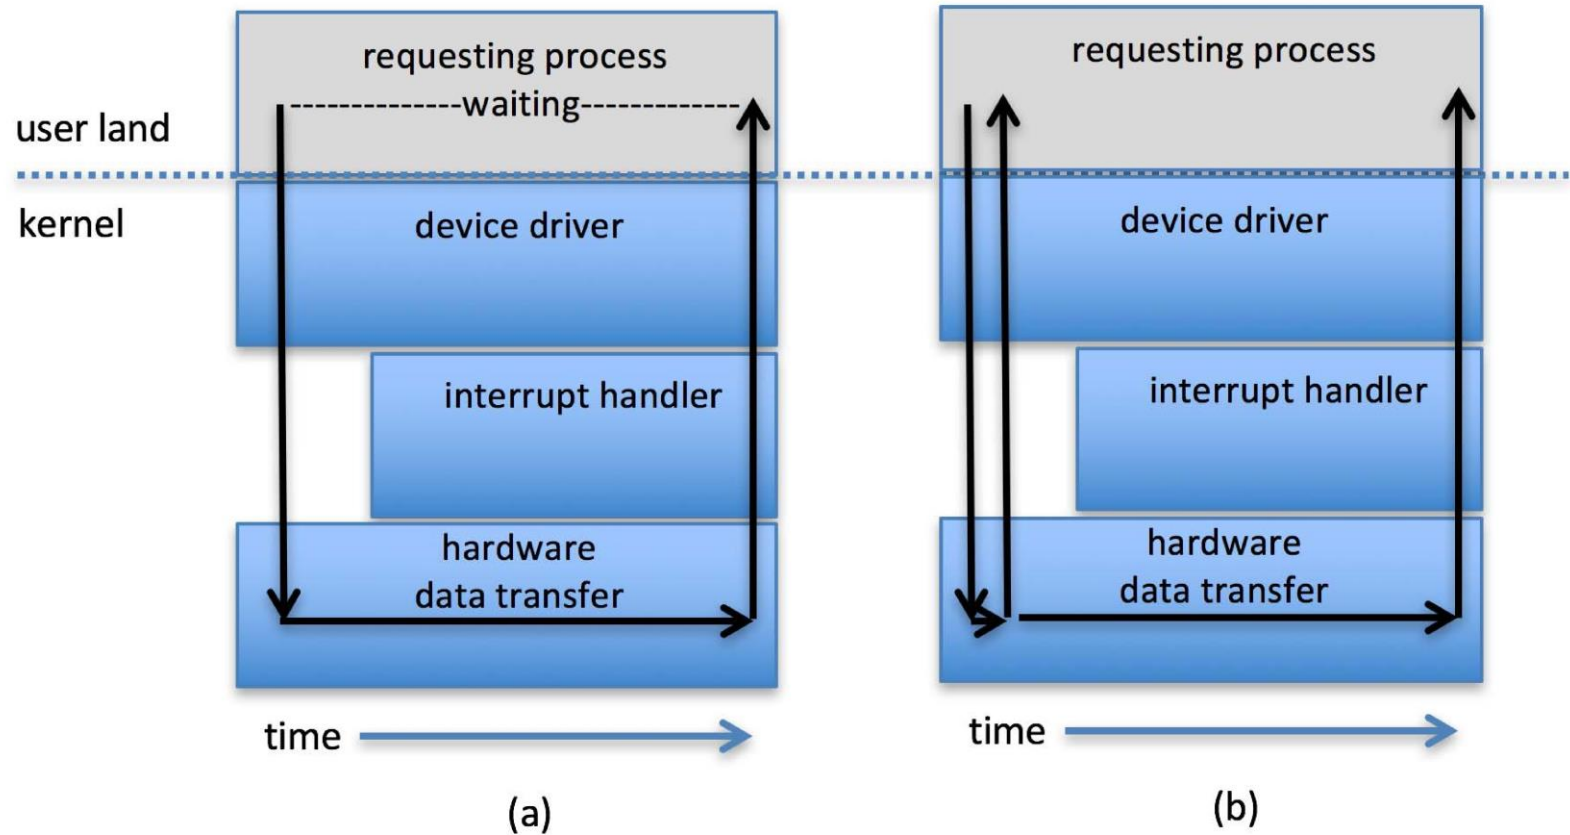
# Blocking and Nonblocking I/O

---

- ▶ **Blocking** - process suspended until I/O completed
  - ▶ Easy to use and understand
  - ▶ Insufficient for some needs
- ▶ **Nonblocking** - I/O call returns as much as available
  - ▶ User interface, data copy (buffered I/O)
  - ▶ Implemented via multi-threading
  - ▶ Returns quickly with count of bytes read or written
- ▶ **Asynchronous** - process runs while I/O executes
  - ▶ Difficult to use
  - ▶ I/O subsystem signals process when I/O completed



# Two I/O Methods





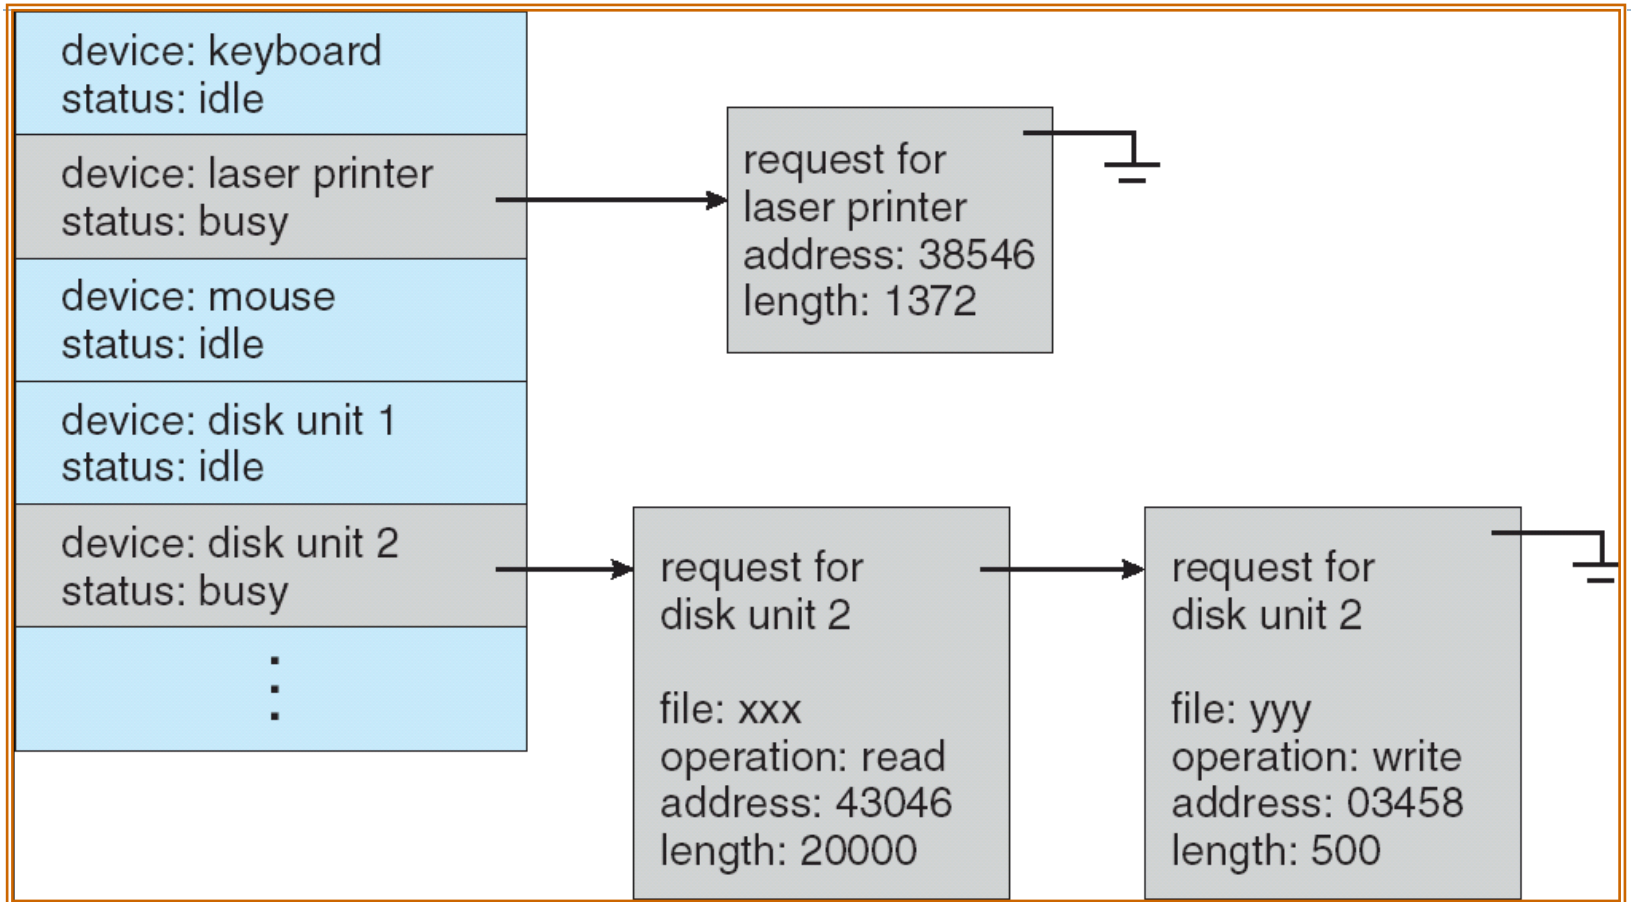
# Kernel I/O Subsystem

---


- ▶ **Scheduling**
  - ▶ Some I/O request ordering via per-device queue
  - ▶ Some OSs try fairness
- ▶ **Buffering - store data in memory while transferring between devices**
  - ▶ To cope with device speed mismatch
  - ▶ To cope with device transfer size mismatch
  - ▶ To maintain “copy semantics”



# Device-status Table



# Common PC and Data-center I/O devices and Interface Speeds



# Kernel I/O Subsystem

---

- ▶ **Caching** - fast memory holding copy of data
  - ▶ Always just a copy
  - ▶ Key to performance
- ▶ **Spooling** - hold output for a device
  - ▶ If device can serve only one request at a time
  - ▶ i.e., Printing
- ▶ **Device reservation** - provides exclusive access to a device
  - ▶ System calls for allocation and deallocation
  - ▶ Watch out for deadlock



# Error Handling

---

- ▶ OS can recover from disk read, device unavailable, transient write failures
- ▶ Most return an error number or code when I/O request fails
- ▶ System error logs hold problem reports



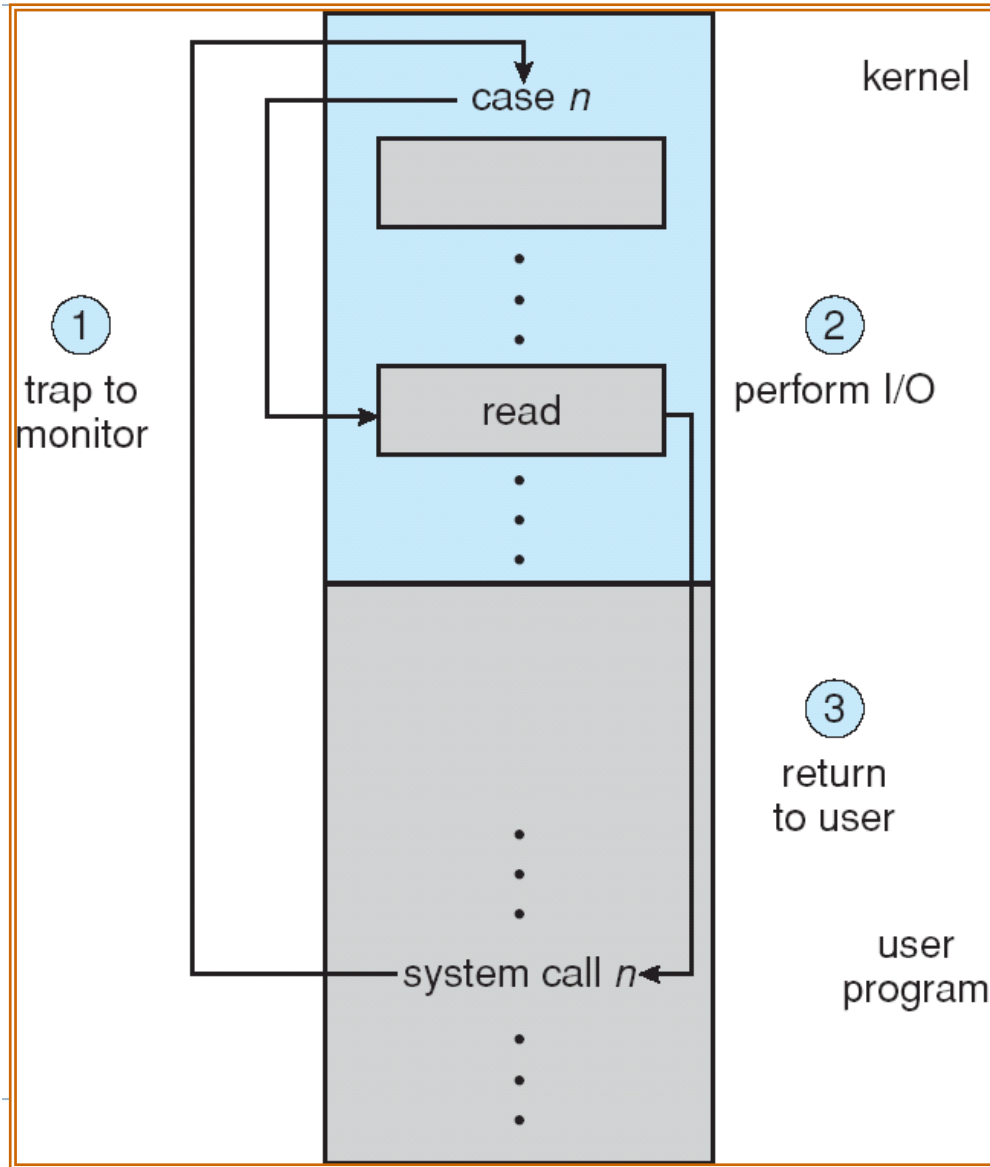
# I/O Protection

---

- ▶ User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions
  - ▶ All I/O instructions defined to be privileged
  - ▶ I/O must be performed via system calls
    - ▶ Memory-mapped and I/O port memory locations must be protected too



# Use of a System Call to Perform I/O



# Kernel Data Structures

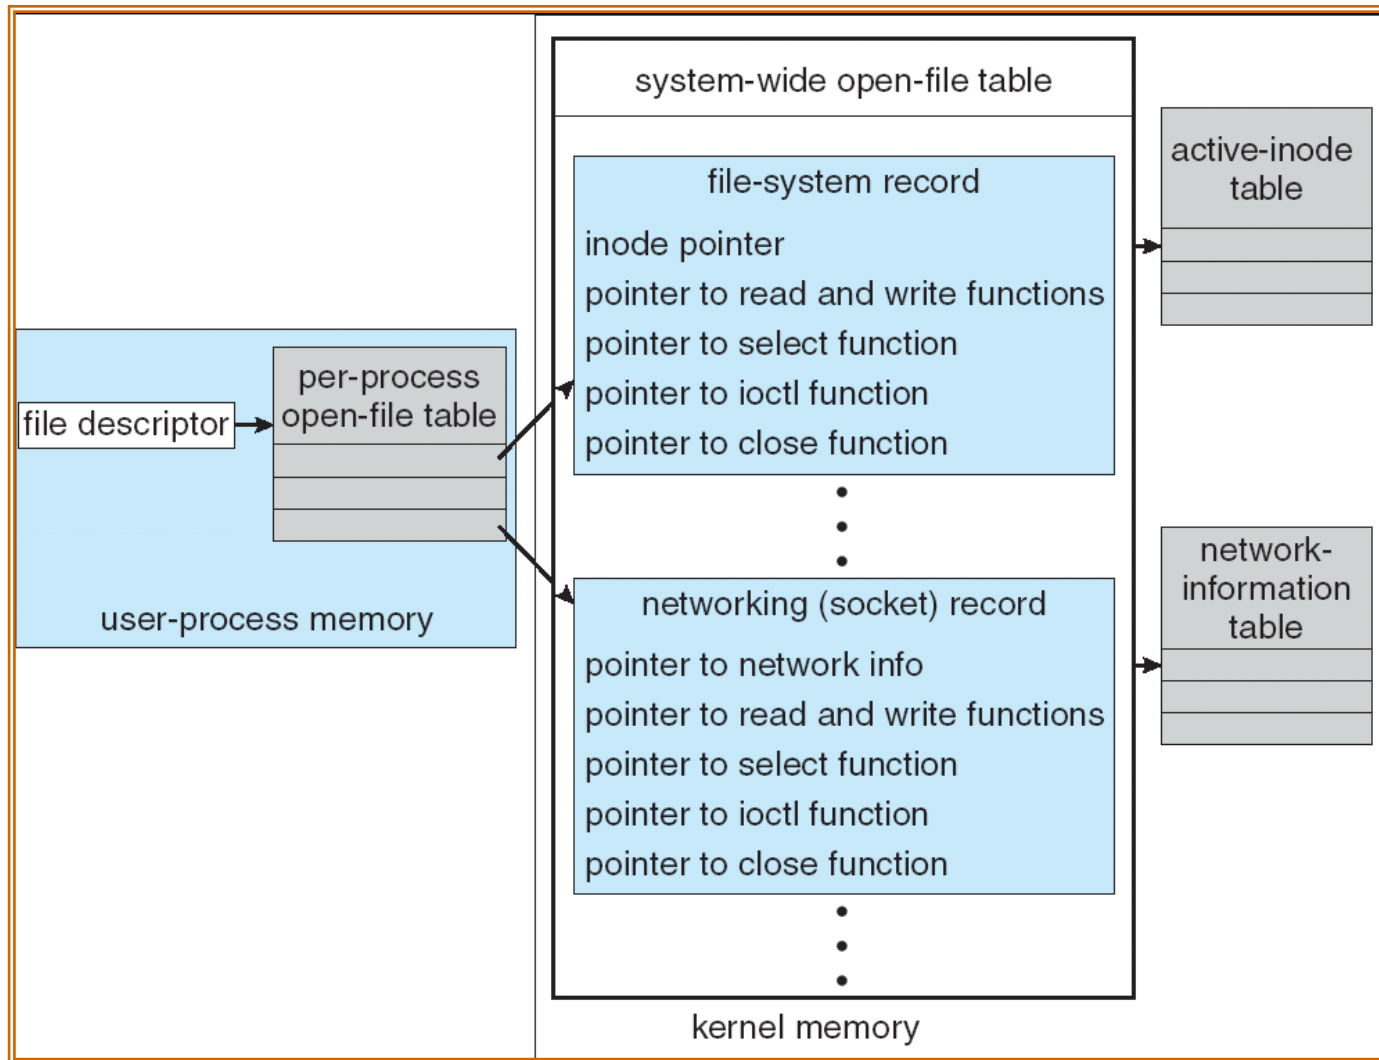
---

- ▶ Kernel keeps state info for I/O components, including open file tables, network connections, character device state
- ▶ Many, many complex data structures to track buffers, memory allocation, “dirty” blocks
- ▶ Some use object-oriented methods and message passing to implement I/O





# UNIX I/O Kernel Structure



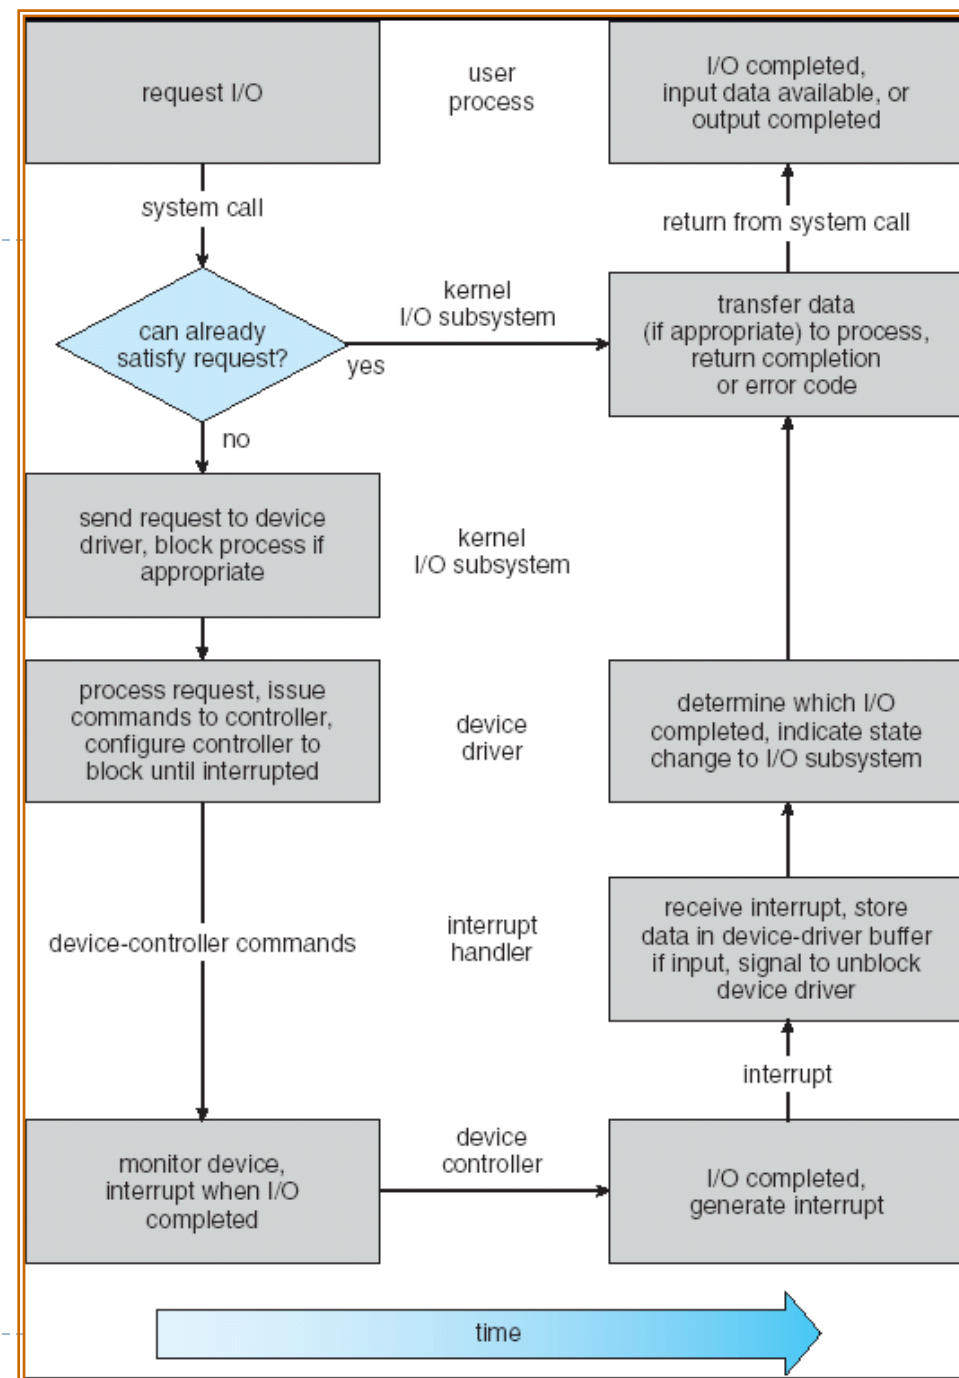
# I/O Requests to Hardware Operations

---

- ▶ Consider reading a file from disk for a process:
  - ▶ Determine device holding file
  - ▶ Translate name to device representation
  - ▶ Physically read data from disk into buffer
  - ▶ Make data available to requesting process
  - ▶ Return control to process



# Life Cycle of An I/O Request



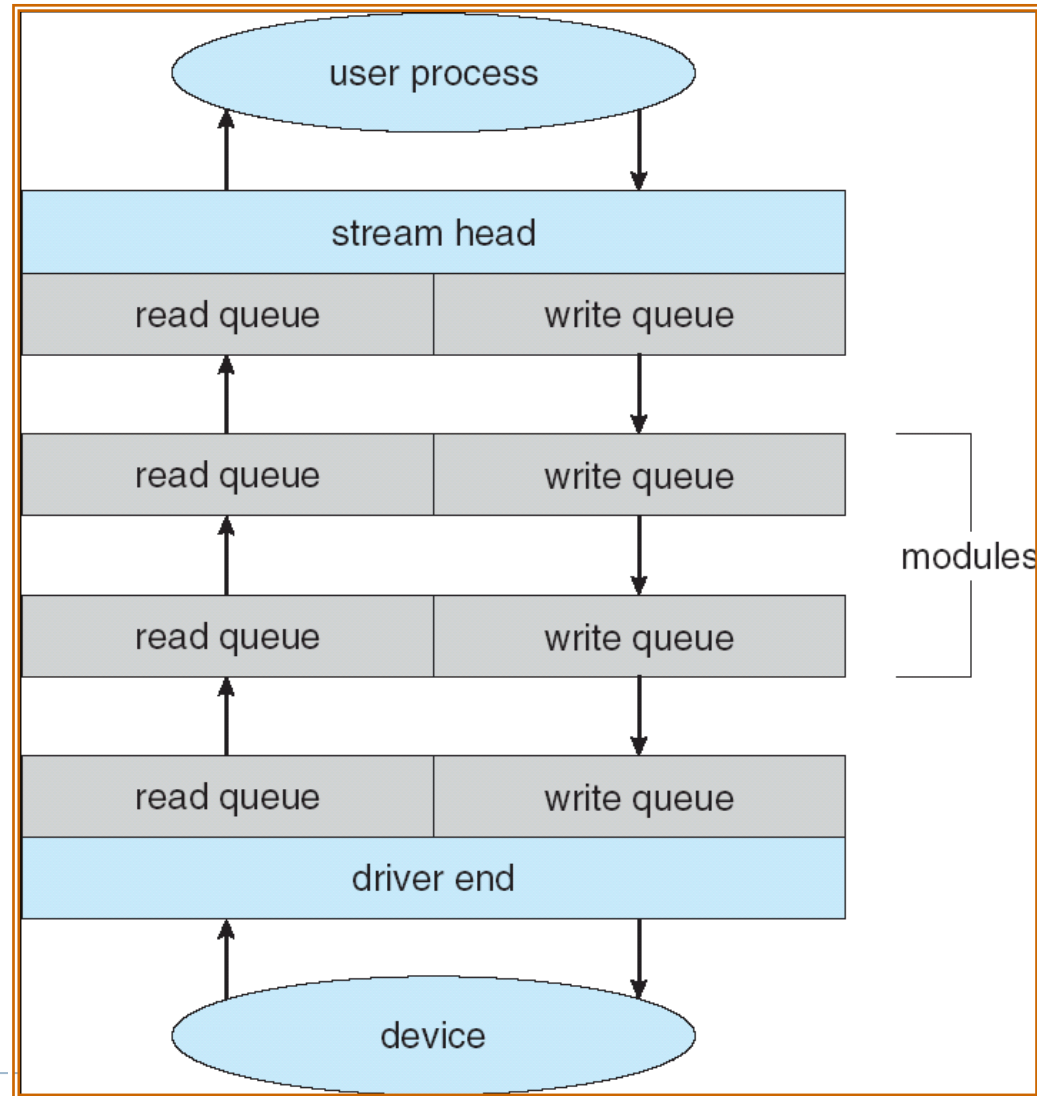
# STREAMS

---

- ▶ **STREAM** – a full-duplex communication channel between a user-level process and a device in Unix System V and beyond
- ▶ A STREAM consists of:
  - STREAM head interfaces with the user process
  - driver end interfaces with the device
  - zero or more STREAM modules between them.
- ▶ Each module contains a **read queue** and a **write queue**
- ▶ Message passing is used to communicate between queues



# The STREAMS Structure



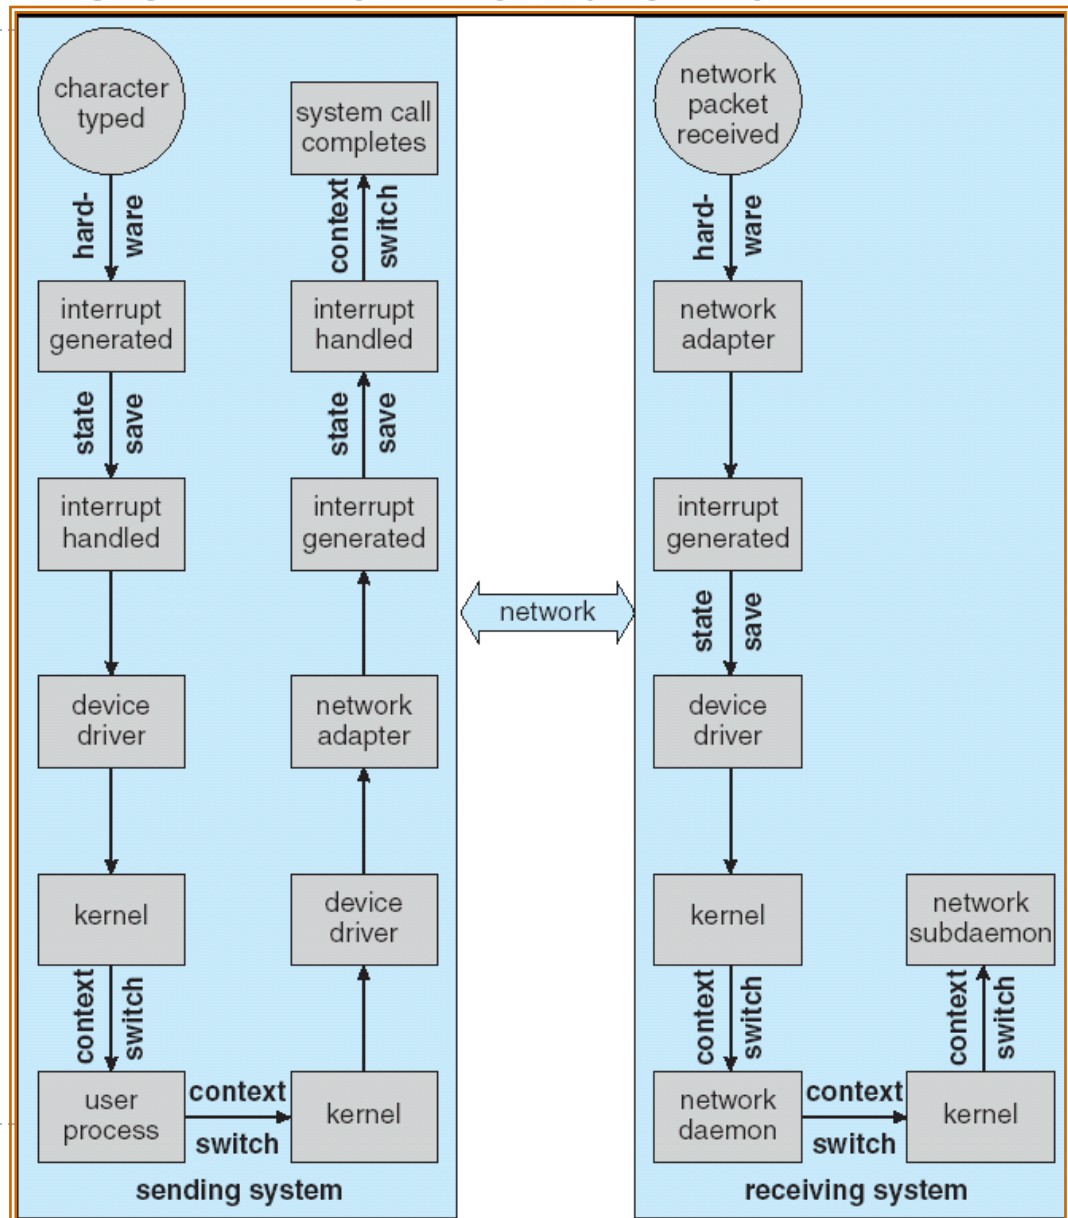
# Performance

---

- ▶ I/O a major factor in system performance:
  - ▶ Demands CPU to execute device driver, kernel I/O code
  - ▶ Context switches due to interrupts
  - ▶ Data copying
  - ▶ Network traffic especially stressful



# Intercomputer Communications



# Improving Performance

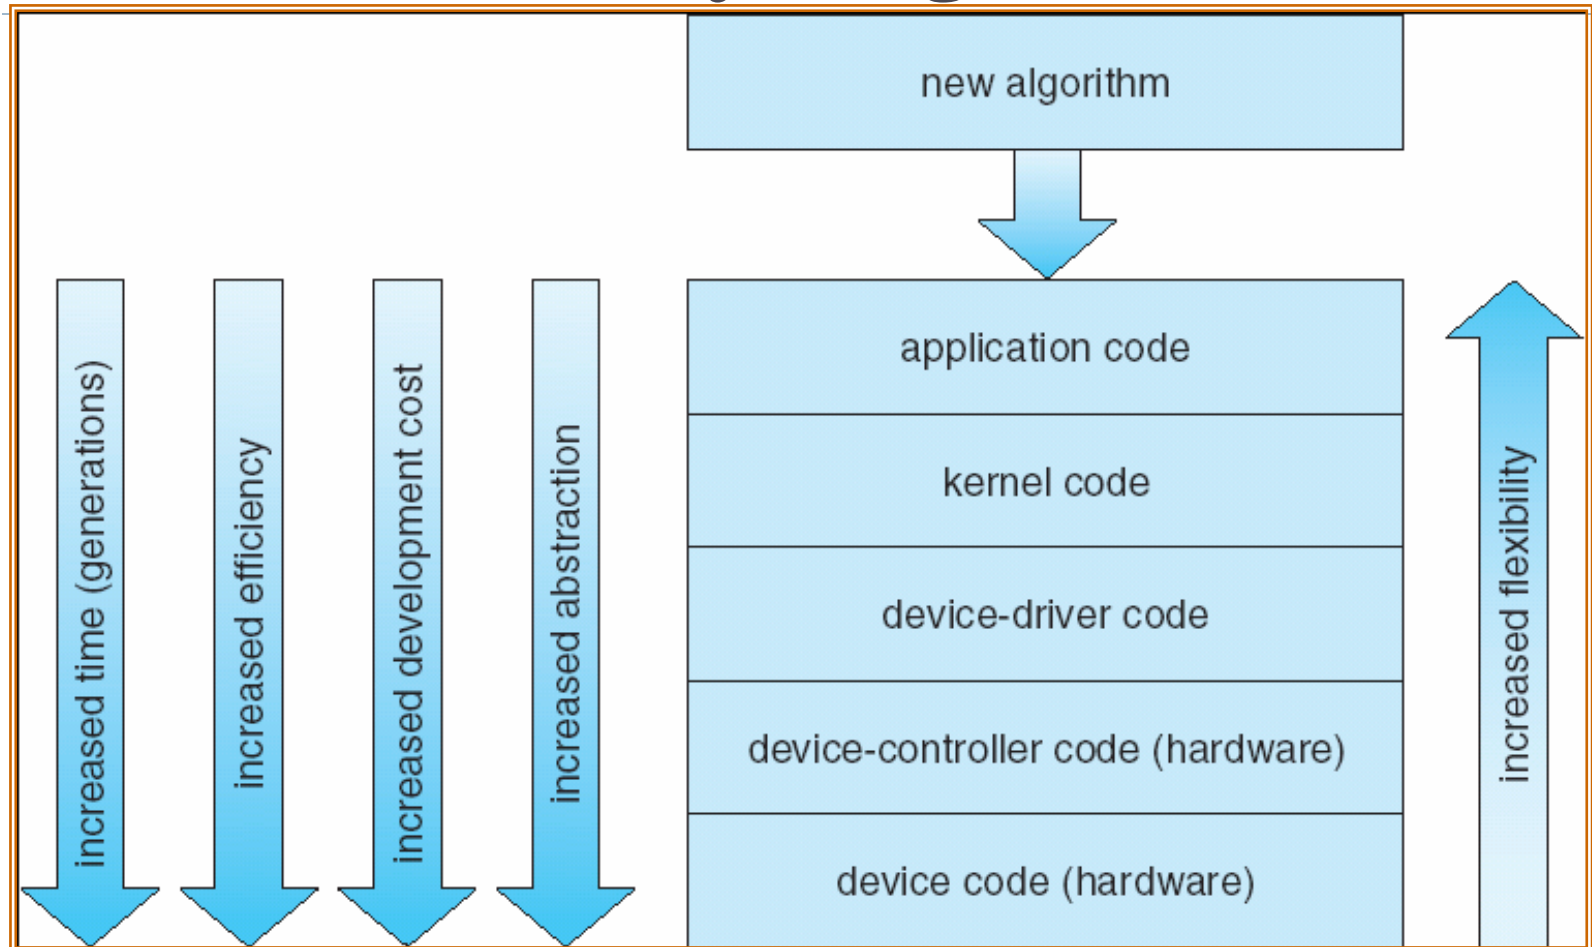
---

- ▶ Reduce number of context switches
- ▶ Reduce data copying
- ▶ Reduce interrupts by using large transfers, smart controllers, polling
- ▶ Use DMA
- ▶ Balance CPU, memory, bus, and I/O performance for highest throughput





# Device-Functionality Progression



# I/O Performance of Storage (and Network Latency)

